
CloudFormation

Command Line Interface

User Guide for Type Development



CloudFormation Command Line Interface: User Guide for Type Development

Copyright © 2020 Amazon Web Services, Inc. and/or its affiliates. All rights reserved.

Amazon's trademarks and trade dress may not be used in connection with any product or service that is not Amazon's, in any manner that is likely to cause confusion among customers, or in any manner that disparages or discredits Amazon. All other trademarks not owned by Amazon are the property of their respective owners, who may or may not be affiliated with, connected to, or sponsored by Amazon.

Table of Contents

What Is the CloudFormation Command Line Interface?	1
Setting Up Your Dev Environment	1
Setting Up Your Environment (macOS)	1
Creating Resource Providers	3
Using the CloudFormation CLI to Create Resource Providers	3
Modeling Resource Providers	4
Defining Property Attributes	5
How to Define a Minimal Resource	5
Resource Provider Schema	5
Developing Resource Providers	14
Implementing Resource Handlers	14
Monitoring Runtime Logging for Resource Types	15
Testing Resource Providers	15
Progress Chaining, Stabilization and Callback Pattern	25
CloudFormation CLI Reference	29
Registering Resource Providers	34
Resource Provider Versions	35
Resource Provider Scope	35
Resource Provider Provisioning	35
Deregistering Resource Providers and Provider Versions	35
Walkthrough: Develop a Provider	36
Prerequisites	36
Create the Resource Provider Development Project	36
Model the Resource Provider	36
Implement the Resource Handlers	39
Test the Resource Provider	55
Submit the Resource Provider	60
Provision the Resource in a CloudFormation Stack	61
Resource Provider FAQ	62
Updates	62
Schema Development	62
Permissions and Authorization	62
Provider Development	63
Testing	63
Deployment	64
Document History	65
AWS glossary	66

What Is the CloudFormation Command Line Interface?

The CloudFormation Command Line Interface (CLI) is an open-source tool that enables you to develop and test AWS and third-party resources, and register them for use in AWS CloudFormation. The CloudFormation CLI provides a consistent way to model and provision both AWS and third-party resources through CloudFormation. It includes commands to enable each step of creating your resource providers.

You can use the CloudFormation CLI to register resource providers—both those you create yourself, as well as ones shared with you—with the CloudFormation registry. This enables you to use CloudFormation capabilities to create, provision, and manage these custom resources in a safe and repeatable manner, just as you would any AWS resource. For more information on the CloudFormation registry, see [Using the CloudFormation Registry](#) in the *CloudFormation User Guide*.

Setting Up Your Environment for Developing Resource Providers

Before you can develop resource providers, you'll need to set up your developer environment, including the CloudFormation CLI.

Currently, plugins are available for the following languages:

- Go
- Java
- Python (Developer Preview)

Or, if you're using another language, you can install the CloudFormation CLI directly.

Setting Up Your Environment (macOS)

Prerequisites

- Python version 3.6 or above
- [AWS Command Line Interface](#) for access to `aws cloudformation` commands.
- Your choice of IDE

The [Walkthrough: Develop a Resource Provider \(p. 36\)](#) walkthrough uses the Community Edition of the [IntelliJ IDEA](#).

- [AWS Serverless Application Model Command Line Interface](#) (AWS SAM CLI)

Note

[Installing the AWS SAM CLI](#) requires Docker as a prerequisite for testing your resource provider locally.

Complete the following steps:

1. Install Homebrew

First, install [Homebrew](#), an open-source package manager for macOS. You'll use Homebrew to install additional development requirements.

2. Next, use Homebrew to install Python and the AWS Command Line Interface (AWS CLI).

```
$ brew update
$ brew install python awscli
```

Installing the CloudFormation CLI and Plugins (macOS)

Use the Python Package Index (PyPI) to install the development plugin for the language of your choice. Installing any of the plugins listed below also installs the CloudFormation CLI. For full installation instructions, refer to the appropriate plugin repository.

Available Language Plugins

Language	Plugin Status	GitHub Location	PyPI Installation
Go	General Availability	cloudformation-cli-go-plugin	cloudformation-cli-go-plugin
Java	General Availability	cloudformation-cli-java-plugin	cloudformation-cli-java-plugin
Python	Developer Preview	cloudformation-cli-python-plugin	N/A

Creating Resource Providers

If you use third-party resources in your infrastructure and applications, you can now model and automate those resources by developing them as *resource providers* for use within CloudFormation. A resource provider includes a resource type specification, and handlers that control API interactions with the underlying AWS or third-party services. These interactions include create, read, update, delete and list (CRUDL) operations for resources. Use resource providers to model and provision resources using CloudFormation.

Resource providers are treated as first-class citizens within CloudFormation; you can use CloudFormation capabilities to create, provision, and manage these custom resources in a safe and repeatable manner, just as you would any AWS resource. Using resource providers for third-party resources provides you a way to reliably manage these resources using a single tool, without having to resort to time-consuming and error-prone methods like manual configuration or custom scripts.

You can create resource providers and make them available for use within the AWS account in which they are registered.

Using the CloudFormation CLI to Create Resource Providers

Use the [AWS CloudFormation Command Line Interface \(CLI\)](#) to develop your resource providers. The CloudFormation CLI is an open-source project that provides a consistent way to model and provision both AWS and third-party resources using CloudFormation. It includes commands to enable each step of creating your resource providers.

There are three major steps in developing a resource providers:

- Model

Create and validate a schema that serves as the canonical definition of your resource provider.

Use the `init` (p. 30) command to generate your resource project, including an example resource schema. Edit the example schema to define the actual model of your resource provider. This includes resource properties and their attributes, as well specifying resource event handlers and any permissions needed for each.

As you iterate on your resource model, you can use the `validate` (p. 31) command to validate your schema against the [Resource Provider Definition Schema](#) and fix any issues.

- Develop

Add logic that controls what happens to the resource at each stage in its lifecycle, and then test the resource locally to ensure it works as expected.

Implement the resource provisioning actions that the CloudFormation CLI stubbed out when you initially generated your resource project.

If you make changes to your resource schema, use the `generate` (p. 31) command to generate the language-specific data model, contract test, and unit test stubs based on the current state of

the resource schema. (If you use the Java add-in for the CloudFormation CLI, this is done for you automatically.)

When you're ready to test the resource behavior, the CloudFormation CLI provides two commands for testing:

- Use the `invoke` command to test a single handler.
- Use the `test` (p. 32) command to run the entire suite of resource contract tests locally, using the AWS SAM Command Line Interface (SAM CLI), to make sure the handlers you've written comply with expected handler behavior at each stage of the resource lifecycle.
- Register

Register the resource provider with the CloudFormation registry in order to make it available for use in CloudFormation templates.

Use the `submit` (p. 33) command to register the resource provider with CloudFormation and make it available for use in CloudFormation operations. Registration includes:

- Validating the resource schema.
- Packaging up the resource project files and uploading them to CloudFormation.
- Registering the resource definition in your account, in the specified region.

You can register multiple versions of a resource provider, and specify which version you want users to use by default.

Modeling Resource Providers for Use in AWS CloudFormation

The first step in creating a resource provider is *modeling* that resource, which involves crafting a schema that defines the resource, its properties, and their attributes. When you initially create your resource provider project using the CloudFormation CLI `init` (p. 30) command, one of the files created is an example resource schema. Use this schema file as a starting point for defining the shape and semantics of your resource provider.

In order to be considered valid, your resource provider's schema must adhere to the [Resource Provider Definition Schema](#). This meta-schema provides a means of validating your resource specification during resource development.

The Resource Provider Definition Schema is a *meta-schema* that extends [draft-07](#) of the [JSON Schema](#). To simplify authoring resource specifications, the Resource Provider Definition Schema constrains the scope of the full JSON Schema standard in terms of how certain validations can be expressed, and encourages consistent modelling for all resource schemas. (For full details on how the Resource Provider Definition Schema differs from the full JSON schema, see [Divergence From JSON Schema](#).)

Once you have defined your resource schema, you can use the CloudFormation CLI `validate` (p. 31) command to verify that the resource schema is valid.

In terms of testing, the resource schema also determines:

- What unit test stubs are generated in your resource package, and what contract tests are appropriate to run for the resource. When you run the CloudFormation CLI `generate` (p. 31) command, the CloudFormation CLI generates empty unit tests based on the properties of the resource and their attributes.
- Which contract tests are appropriate for CloudFormation CLI to run for your resources. When you run the `test` (p. 32) command, the CloudFormation CLI runs the appropriate contract tests, based on which handlers are included in your resource schema.

Defining Property Attributes

Certain properties of a resource may have special meaning when used in different contexts. For example, a given resource property may be read-only when read back for state changes, but can be specified when used as the target of a `$ref` from a related resource. Because of this semantic difference in how this property metadata should be interpreted, certain property attributes are defined at the resource level, rather than at a property level.

These attributes include:

- `primaryIdentifier`
- `additionalIdentifiers`
- `createOnlyProperties`
- `readOnlyProperties`
- `writeOnlyProperties`

For reference information on resource schema elements, see [Resource Provider Schema \(p. 5\)](#).

How to Define a Minimal Resource

The example below displays a minimal resource provider definition. In this case, the resource consists of a single optional property, `Name`, which is also specified as its primary (and only) identifier.

Note that this resource schema would require a `handlers` section with the create, read, and update handlers specified in order for the resource to actually be provisioned within a CloudFormation account.

```
{
  "typeName": "myORG::myService::myResource",
  "properties": {
    "Name": {
      "description": "The name of the resource.",
      "type": "string",
      "pattern": "^[a-zA-Z0-9_-]{0,64}$",
      "maxLength": 64
    }
  },
  "createOnlyProperties": [
    "/properties/Name"
  ],
  "identifiers": [
    [
      "/properties/Name"
    ]
  ],
  "additionalProperties": false
}
```

Resource Provider Schema

In order to be considered valid, your resource provider's schema must adhere to the [Resource Provider Definition Schema](#). This meta-schema provides a means of validating your resource specification during resource development.

Syntax

Below is the structure for a typical resource provider schema. For the complete meta-schema definition, see the [Resource Provider Definition Schema](#) on [GitHub](#).


```
{
  "typeName": "string",
  "description": "string",
  "sourceUrl": "string",
  "documentationUrl": "string",
  "definitions": {
    "definitionName": {
      . . .
    }
  },
  "properties": {
    "propertyName": {
      "description": "string",
      "type": "string",
      . . .
    },
  },
  "required": [
    "propertyName"
  ],
  "readOnlyProperties": [
    "/properties/propertyName"
  ],
  "writeOnlyProperties": [
    "/properties/propertyName"
  ],
  "primaryIdentifier": [
    "/properties/propertyName"
  ],
  "createOnlyProperties": [
    "/properties/propertyName"
  ],
  "deprecatedProperties": [
    "/properties/propertyName"
  ],
  "additionalIdentifiers": [
    [
      "/properties/propertyName"
    ]
  ],
  "handlers": {
    "create": {
      "permissions": [
        "permission"
      ]
    },
    "read": {
      "permissions": [
        "permission"
      ]
    },
    "update": {
      "permissions": [
        "permission"
      ]
    },
    "delete": {
      "permissions": [
        "permission"
      ]
    },
    "list": {
      "permissions": [
```

```
    "permission"  
  ]  
}  
}  
}
```

Properties

typeName

The unique name for your resource. Specifies a three-part namespace for your resource, with a recommended pattern of `Organization::Service::Resource`.

Note

The following organization namespaces are reserved and cannot be used in your resource type names:

- Alexa
- AMZN
- Amazon
- ASK
- AWS
- Custom
- Dev

Required: Yes

Pattern: `^[a-zA-Z0-9]{2,64}:[a-zA-Z0-9]{2,64}:[a-zA-Z0-9]{2,64}$`

description

A short description of the resource. This will be displayed in the AWS CloudFormation console.

Required: Yes

sourceUrl

The URL of the source code for this resource, if public.

documentationUrl

The URL of a page providing detailed documentation for this resource.

Note

While the resource schema itself should include complete and accurate property descriptions, the `documentationUrl` property enables you to provide users with documentation that describes and explains the resource in more detail, including examples, use cases, and other detailed information.

definitions

Use the `definitions` block to provide shared resource property schemas.

It is considered a best practice is to use the `definitions` section to define schema elements that may be used at multiple points in your resource provider schema. You can then use a JSON pointer to reference that element at the appropriate places in your resource provider schema.

properties

The properties of the resource.

All properties of a resource must be expressed in the schema. Arbitrary inputs are not allowed. A resource must contain at least one property.

Nested properties are not allowed. Instead, define any nested properties in the `definitions` element, and use a `$ref` pointer to reference them in the desired property.

Required: Yes

propertyName

insertionOrder

For properties of type `array`, set to `true` to specify that the order in which array items are specified must be honored, and that changing the order of the array will indicate a change in the property.

The default is `false`.

readOnly

Reserved for CloudFormation use.

writeOnly

Reserved for CloudFormation use.

dependencies

Any properties that are required if this property is specified.

patternProperties

Use to specify a specification for key-value pairs.

```
"type": "object",  
"propertyNames": {  
  "format": "regex"  
}
```

properties

Minimum: 1

patternProperties

Pattern: `^[A-Za-z0-9]{1,64}$`

Specifies a pattern that properties must match to be valid.

allOf

The property must contain all of the data structures define here.

Contains a single schema. A list of schemas is not allowed.

Minimum: 1

anyOf

The property can contain any number of the data structures define here.

Contains a single schema. A list of schemas is not allowed.

Minimum: 1

oneOf

The property must contain only one of the data structures define here.

Contains a single schema. A list of schemas is not allowed.

Minimum: 1

items

For properties of type array, defines the data structure of each array item.

Contains a single schema. A list of schemas is not allowed.

In addition, the following elements, defined in [draft-07](#) of the [JSON Schema](#), are allowed in the `properties` object:

- `$ref`
- `$comment`
- `title`
- `description`
- `examples`
- `default`
- `multipleOf`
- `maximum`
- `exclusiveMaximum`
- `minimum`
- `exclusiveMinimum`
- `exclusiveMinimum`
- `minLength`
- `pattern`
- `maxItems`
- `minItems`
- `uniqueItems`
- `contains`
- `maxProperties`
- `maxProperties`
- `required`
- `const`
- `enum`
- `type`
- `format`

remote

Reserved for CloudFormation use.

readOnlyProperties

Resource properties that can be returned by a `read` or `list` request, but cannot be set by the user.

Type: List of JSON pointers

`writeOnlyProperties`

Resource properties that can be specified by the user, but cannot be returned by a `read` or `list` request. Write-only properties are often used to contain passwords, secrets, or other sensitive data.

Type: List of JSON pointers

`createOnlyProperties`

Resource properties that can be specified by the user only during resource creation.

Note

Any property not explicitly listed in the `createOnlyProperties` element can be specified by the user during a resource update operation.

Type: List of JSON pointers

`deprecatedProperties`

Resource properties that have been deprecated by the underlying service provider. These properties are still accepted in create and update operations. However they may be ignored, or converted to a consistent model on application. Deprecated properties are not guaranteed to be returned by read operations.

Type: List of JSON pointers

`primaryIdentifier`

The uniquely identifier for an instance of this resource provider. An identifier is a non-zero-length list of JSON pointers to properties that form a single key. An identifier can be a single or multiple properties to support composite-key identifiers.

Type: List of JSON pointers

Required: Yes

`additionalIdentifiers`

An optional list of supplementary identifiers, each of which uniquely identifies an instance of this resource provider. An identifier is a non-zero-length list of JSON pointers to properties that form a single key. An identifier can be a single property, or multiple properties to construct composite-key identifiers.

Type: List of JSON pointers

Minimum: 1

`handlers`

Specifies the provisioning operations which can be performed on this resource provider. The handlers specified determine what provisioning actions CloudFormation takes with respect to the resource during various stack operations.

- If the resource provider does not contain `create`, `read`, and `delete` handlers, CloudFormation cannot actually provision the resource.
- If the resource provider does not contain an `update` handler, CloudFormation cannot update the resource during stack update operations, and will instead replace it.

If your resource type calls AWS APIs in any of its handlers, you must create an [IAM execution role](#) that includes the necessary permissions to call those AWS APIs, and provision that execution role in your account. For more information, see [Accessing AWS APIs from a Resource Type](#).

`create`

`permissions`

A string array that specifies the AWS permissions needed to invoke the create handler.

You must specify at least one permission for each handler.

Required: Yes

read

permissions

A string array that specifies the AWS permissions needed to invoke the read handler.

You must specify at least one permission for each handler.

Required: Yes

update

permissions

A string array that specifies the AWS permissions needed to invoke the update handler.

You must specify at least one permission for each handler.

Required: Yes

delete

permissions

A string array that specifies the AWS permissions needed to invoke the delete handler.

You must specify at least one permission for each handler.

Required: Yes

list

The `list` handler must at least return the resource's primary identifier.

permissions

A string array that specifies the AWS permissions needed to invoke the list handler.

You must specify at least one permission for each handler.

Required: Yes

allOf

The resource must contain all of the data structures defined here.

Minimum: 1

anyOf

The resource can contain any number of the data structures define here.

Minimum: 1

oneOf

The resource must contain only one of the data structures define here.

Minimum: 1

In addition, the following element, defined in [draft-07](#) of the [JSON Schema](#), is allowed:

- required

How to Specify a Property as Dependent on Another

Use the `dependencies` element to specify if a property is required in order for another property to be specified. In the following example, if the user specifies a value for the `ResponseCode` property, they must also specify a value for `ResponsePagePath`, and vice versa. (Note that, as a best practice, this is also called out in the description of each property.)

```
"properties": {
  "CustomErrorResponse": {
    "additionalProperties": false,
    "dependencies": {
      "ResponseCode": [
        "ResponsePagePath"
      ],
      "ResponsePagePath": [
        "ResponseCode"
      ]
    },
    "properties": {
      "ResponseCode": {
        "description": "The HTTP status code that you want CloudFront to return to the viewer along with the custom error page. If you specify a value for ResponseCode, you must also specify a value for ResponsePagePath.",
        "type": "integer"
      },
      "ResponsePagePath": {
        "description": "The path to the custom error page that you want CloudFront to return to a viewer when your origin returns the HTTP status code specified by ErrorCode. If you specify a value for ResponsePagePath, you must also specify a value for ResponseCode.",
        "type": "string"
      }
      . . .
    },
    "type": "object"
  },
  . . .
}
```

How to Define Nested Properties

It is considered a best practice is to use the `definitions` section to define schema elements that may be used at multiple points in your resource provider schema. You can then use a JSON pointer to reference that element at the appropriate places in your resource provider schema.

For example, define the reused element in the `definitions` section:

```
"definitions": {
  "AccountId": {
    "pattern": "^[0-9]{12}$",
    "type": "string"
  },
  . . .
},
. . .
```

And then reference that definition where appropriate:

```
"AwsAccountNumber": {
  "description": "An AWS account that is included in the TrustedSigners complex type for
this distribution.",
  "$ref": "#/definitions/AccountId"
},
. . .
```

Advanced: How to Encapsulate Complex Logic

Use the `allof`, `oneOf`, or `anyOf` elements to encapsulate complex logic in your resource provider schema.

In the example below, if `whitelist` is specified for the `Forward` property in your resource, then the `WhitelistedNames` property must also be specified.

```
"properties": {
  "Cookies": {
    "oneOf": [
      {
        "additionalProperties": false,
        "properties": {
          "Forward": {
            "description": "Specifies which cookies to forward to the origin for
this cache behavior.",
            "enum": [
              "all",
              "none"
            ],
            "type": "string"
          }
        },
        "required": [
          "Forward"
        ]
      },
      {
        "additionalProperties": false,
        "properties": {
          "Forward": {
            "description": "Specifies which cookies to forward to the origin for
this cache behavior.",
            "enum": [
              "whitelist"
            ],
            "type": "string"
          },
          "WhitelistedNames": {
            "description": "Required if you specify whitelist for the value of
Forward.",
            "items": {
              "type": "string"
            },
            "minItems": 1,
            "type": "array"
          }
        },
        "required": [
          "Forward",
          "WhitelistedNames"
        ]
      }
    ],
    "type": "object"
  }
}
```



```
} ,  
} ,  
. . .
```

Developing Resource Providers for Use in AWS CloudFormation Templates

Once you've modeled your resource provider, and validated its schema, the next step is to develop the resource. Developing the resource consists of two main steps:

- Implementing the appropriate event handlers for your resource.
- Testing the resource locally to ensure it works as expected.

Implementing Resource Handlers

When you [generate](#) (p. 31) your resource package, the CloudFormation CLI stubs out empty handler functions, each of which corresponds to a specific event in the resource lifecycle. You add logic to these handlers to control what happens to your resource provider at each stage of its lifecycle.

- `create`: CloudFormation invokes this handler when the resource is initially created during stack create operations.
- `read`: CloudFormation invokes this handler as part of a stack update operation when detailed information about the resource's current state is required.
- `update`: CloudFormation invokes this handler when the resource is updated as part of a stack update operation.
- `delete`: CloudFormation invokes this handler when the resource is deleted, either when the resource is deleted from the stack as part of a stack update operation, or the stack itself is deleted.
- `list`: CloudFormation invokes this handler when summary information about multiple resources of this resource provider is required.

You can only specify a single handler for each event.

Which handlers you implement for a resource determine what provisioning actions CloudFormation takes with respect to the resource during various stack operations:

- If the resource provider contains both `create` and `update` handlers, CloudFormation invokes the appropriate handler during stack create and update operation.
- If the resource does not contain an `update` handler, CloudFormation cannot update the resource during stack update operations, and will instead replace it. CloudFormation invokes the `create` handler to create a new resource, then deletes the old resource by invoking the `delete` handler.
- If the resource provider does not contain `create`, `read`, and `delete` handlers, CloudFormation cannot actually provision the resource.

Use the resource schema to specify which handlers you have implemented. If you choose not to implement a specific handler, remove it from the `handlers` section of the resource schema.

Accessing AWS APIs from a Resource Type

If your resource type calls AWS APIs in any of its handlers, you must create an [IAM execution role](#) that includes the necessary permissions to call those AWS APIs, and provision that execution role in your

account. CloudFormation then assumes that execution role to provide your resource type with the appropriate credentials.

When you call `generate`, the CloudFormation CLI automatically generates an execution role template, `resource-role.yaml`, as part of generating the code files for the resource type package. This template is based on the permissions specified for each handler in the [handlers](#) section of the [Resource Provider Schema](#). When you use `submit` to register the resource type, the CloudFormation CLI attempts to create or update an execution role based on the template, and then passes this execution role to CloudFormation as part of the registration.

For more information on the permissions available per AWS service, see [Actions, Resources, and Condition Keys for AWS Services](#) in the *AWS Identity and Access Management User Guide*.

Monitoring Runtime Logging for Resource Types

When you register a resource type using `cfn submit`, CloudFormation creates a CloudWatch log group for the resource type in your account. This enables you to access the logs for your resource to help you diagnose any faults. The log group is named according to the following pattern:

```
/my-resource-type-stack-ResourceHandler-string
```

Where:

- `my-resource-type` is the three-part resource type name
- `string` is a unique string generated by CloudFormation

Now, when you execute stack operations for stacks that contain the resource type, CloudFormation delivers log events emitted by the resource type to this log group.

Testing Resource Providers Using Contract Tests

As you model and develop your resource provider, you can have the CloudFormation CLI perform tests to ensure the resource provider is behaving as expected during each event in the resource lifecycle. The CloudFormation CLI performs a suite of tests, each written to test a requirement contained in the resource provider handler contract.

Testing Resource Types Locally Using SAM

Once you've implemented the desired handlers for your resource, you can also test the resource locally using the AWS SAM command line interface (CLI), to make sure your resource behaves as expected, debug what's wrong, and fix any issues.

To start testing, use the SAM CLI to start the Local Lambda Service. Run the following command in a terminal separate from your resource provider workspace, or as a background process.

```
$ sam local start-lambda
```

If you have functions defined in your SAM template, it will provide an endpoint to invoke these functions locally. This is especially helpful because it allows for remote debugging to step through resource provider invocations in real time.

```
Starting the Local Lambda Service. You can now invoke your Lambda Functions defined in your
template through the endpoint.
2020-01-15 15:27:19 * Running on http://127.0.0.1:3001/ (Press CTRL+C to quit)
```

Alternatively, you can also specify using the public Lambda Service and invoke functions deployed in your account. Be aware, however, that using the local service allows for more iteration. To specify a debug port for remote debugging, use the `-d` option:

```
sam local start-lambda -d PORT
```

Once you have the Lambda service started, use the `test` command to perform contract tests:

```
cfn test
```

The CloudFormation CLI selects the appropriate contract tests to execute, based on the handlers specified in your resource provider schema. If a test fail, the CloudFormation CLI outputs a detailed trace of the failure, including the related assertion failure and mismatched values.

For more information on testing using AWS SAM CLI, see [Testing and Debugging Serverless Applications](#) in the *AWS Serverless Application Model Developer Guide*.

How the CloudFormation CLI Constructs and Executes Contract Tests

The CloudFormation CLI uses [PyTest](#), an open-source testing framework, to execute the contract tests.

The tests themselves are located in the `suite` folder of the CloudFormation CLI repository on GitHub. Each test is adorned with the appropriate `pytest` markers. For example, tests applicable to the `create` handler are adorned with the `@pytest.mark.create` marker. This enables the CloudFormation CLI to execute only those tests appropriate for a resource provider, based on the handlers specified in the resource provider's schema. For example, suppose a resource provider's schema specified `create`, `read`, and `delete` handlers. In this case, the CloudFormation CLI would not perform any test marked with only the `@pytest.mark.update` or `@pytest.mark.list`, since those handlers were not implemented.

To test `create` and `update` handlers, the CloudFormation CLI uses the resource provider's schema and [Hypothesis](#), an open-source Python library used to generate testing strategies. The resource provider schema is walked to generate a strategy for valid resource models, and the strategy is used to generate models for tests.

Tests `create`, `update`, and `delete` resources to test various aspects of the resource handler contract during handler operations. The CloudFormation CLI uses `PyTest fixtures` to decrease the amount of time the contract tests take to perform. Using fixtures enable the tests within a test module to share resources, rather than have to create a new resource for each test. Currently, the contract tests employ the following fixtures:

- `created_resource` in the `handler_create` test module
- `updated_resource` in the `handler_update` test module
- `deleted_resource` in `handler_delete` test module

Resource Provider Handler Contracts

The resource provider handler contract specifies the expected and required behavior to which a resource must adhere in each given event handler. It defines a set of specific, unambiguous rules with which `create`, `read`, `update`, `delete` and `list` resource handlers must comply. Following the contract will allow customers to interact with all resource providers under a uniform set of behaviors and expectations, and prevents creation of unintended or duplicate resources.

A resource implementation **MUST** pass all resource contract tests in order to be registered.

Assuming no other concurrent interaction on the resource, the handlers **MUST** comply with the following contract.

All terminology in the handler contract requirements adheres to the [RFC 2119 specification](#).

Create Handlers

- A `create` handler **MUST** always return a `ProgressEvent` object within 60 seconds. For more information, see [ProgressEvent Object Schema](#).

In every `ProgressEvent` object, the `create` handler **MUST** return a model which conforms to the shape of the resource schema. For more information, see [Returned models must conform to the shape of the schema](#).

Every model **MUST** include the `primaryIdentifier`. The only exception is if the first progress event is `FAILED`, and the resource has not yet been created. In this case, a subsequent `read` call **MUST** return `NotFound`.

- A `create` handler **MUST NOT** return `SUCCESS` until it has applied all properties included in the `create` request. For more information, see [Update, create, and delete handlers must satisfy desired-state stabilization](#).

- A `create` handler **MUST** return `IN_PROGRESS` if it has not yet reached the desired-state.

A `create` handler **SHOULD** return a model containing all properties set so far and nothing more during each `IN_PROGRESS` event.

- A `create` handler **MUST** return `FAILED` progress event if it cannot reach the desired-state within the timeout specified in the resource schema.

The progress event **MUST** return an error message and the most applicable error code. For more information, see [Handler error codes](#).

- A `create` handler **MAY** return `SUCCESS` once it reaches the desired-state.

Once the desired state has been reached, a `create` handler **MAY** perform runtime-state stabilization. For more information, see [Update and create handlers should satisfy runtime-state stabilization](#).

When the `create` handler returns `SUCCESS`, it **MUST** return a `ProgressEvent` object containing a model that satisfies the following requirements:

- All properties specified in the `create` request **MUST** be present in the model returned, and they **MUST** match exactly, with the exception of properties defined as `writeOnlyProperties` in the resource schema.
- The model **MUST NOT** return properties defined as `writeOnlyProperties` in the resource schema.
- The model **MUST** contain all properties that have values, including any properties that have default values, and any `readOnlyProperties` as defined in the resource schema.
- The model **MUST NOT** return any properties that are null or do not have values.
- After a `create` operation returns `SUCCESS`, a subsequent `read` request **MUST** succeed when passed in the `primaryIdentifier` or any `additionalIdentifiers` associated with the provisioned resource instance.
- After a `create` operation returns `SUCCESS`, a subsequent `list` operation **MUST** return the `primaryIdentifier` associated with the provisioned resource instance.

If the `list` operation is paginated, the entire `list` operation is defined as all `list` requests until the `nextToken` is null.

- A `create` handler **MUST** be idempotent. A `create` handler **MUST NOT** create multiple resources given the same idempotency token.
- A `create` handler **MUST** return `FAILED` with an `AlreadyExists` error code if the resource already existed prior to the `create` request.

Update Handlers

- An update handler MUST always return a `ProgressEvent` object within 60 seconds. For more information, see [ProgressEvent Object Schema](#).

In every `ProgressEvent` object, the update handler MUST return a model which conforms to the shape of the resource schema. For more information, see [Returned models must conform to the shape of the schema](#).

Every model MUST include the `primaryIdentifier`.

- An update handler MUST NOT return `SUCCESS` until it has applied all properties included in the update request. For more information, see [Update, create, and delete handlers must satisfy desired-state stabilization](#).
- An update handler MUST return `IN_PROGRESS` if it has not yet reached the desired-state.

An update handler SHOULD return a model containing all properties set so far and nothing more during each `IN_PROGRESS` event.

- An update handler MUST return `FAILED` progress event if it cannot reach the desired-state within the timeout specified in the resource schema.

The progress event MUST return an error message and the most applicable error code. For more information, see [Handler error codes](#).

- An update handler MAY return `SUCCESS` once it reaches the desired-state.

Once the desired state has been reached, an update handler MAY perform runtime-state stabilization. For more information, see [Update and create handlers should satisfy runtime-state stabilization](#).

When the update handler returns `SUCCESS`, it MUST return a `ProgressEvent` object containing a model that satisfies the following requirements:

- All properties specified in the update request MUST be present in the model returned, and they MUST match exactly, with the exception of properties defined as `writeOnlyProperties` in the resource schema.
- The model MUST NOT return properties defined as `writeOnlyProperties` in the resource schema.
- The model MUST contain all properties that have values, including any properties that have default values, and any `readOnlyProperties` as defined in the resource schema.
- The model MUST NOT return any properties that are null or do not have values.

All list or collection properties MUST be applied in full. The successful outcome MUST be replacement of the previous properties, if any.

- An update handler MUST return `FAILED` with a `NotUpdatable` error code if the difference between the previous state and the properties included in the request contains a property defined in the `createOnlyProperties` in the resource schema.
- An update handler MUST return `FAILED` with a `NotFound` error code if the resource did not exist prior to the update request.
- An update handler MUST NOT create a new physical resource.

The `primaryIdentifier` returned in every progress event must match the `primaryIdentifier` passed into the request.

Delete Handlers

- A delete handler MUST always return a `ProgressEvent` object within 60 seconds. For more information, see [ProgressEvent Object Schema](#).

- A `delete` handler MUST NOT return `SUCCESS` until the resource has reached the desired state for deletion. For more information, see [Update, create, and delete handlers must satisfy desired-state stabilization](#).
- A `delete` handler MUST return `IN_PROGRESS` if it has not yet reached the desired state.
- A `delete` handler MUST return `FAILED` progress event if it cannot reach the desired-state within the timeout specified in the resource schema.

The progress event MUST return an error message and the most applicable error code. For more information, see [Handler error codes](#).

- A `delete` handler MUST return `SUCCESS` once it reaches the desired state. (This is because there is no runtime-state stabilization for delete requests.)

When the `delete` handler returns `SUCCESS`, the `ProgressEvent` object MUST NOT contain a model.

- A `delete` handler MUST return `FAILED` with a `NotFound` error code if the resource did not exist prior to the delete request.
- Once a `delete` operation successfully completes, any subsequent `update`, `delete`, or `read` request for the deleted resource instance MUST return `FAILED` with a `NotFound` error code.
- Once a `delete` operation successfully completes, any subsequent `list` operation MUST NOT return the `primaryIdentifier` associated with the deleted resource instance.

If the `list` operation is paginated, the 'list operation' is defined as all `list` calls until the `nextToken` is `null`.

- Once a `delete` operation successfully completes, a subsequent `create` request with the same `primaryIdentifier` or `additionalIdentifiers` MUST NOT return `FAILED` with an `AlreadyExists` error code.
- Once a `delete` operation successfully completes, the resource SHOULD NOT be billable to the customer.

Read Handlers

- A `read` handler MUST always return a `ProgressEvent` object within 30 seconds. For more information, see [ProgressEvent Object Schema](#).

A `read` handler MUST always return a status of `SUCCESS` or `FAILED`; it MUST NOT return a status of `IN_PROGRESS`.

- A `read` handler MUST return a model representation that conforms to the shape of the resource schema.
 - The model MUST contain all properties that have values, including any properties that have default values and any `readOnlyProperties` as defined in the resource schema.
 - The model MUST NOT return any properties that are null or do not have values.
 - The model MUST NOT return properties defined as `writeOnlyProperties` in the resource schema.
- A `read` handler MUST return `FAILED` with a `NotFound` error code if the resource does not exist.

List Handlers

- A `list` handler MUST always return a `ProgressEvent` object within 30 seconds. For more information, see [ProgressEvent Object Schema](#).

A `list` handler MUST always return a status of `SUCCESS` or `FAILED`; it MUST NOT return a status of `IN_PROGRESS`.

- A `list` handler MUST return an array of primary identifiers.

When passed in a read request, each `primaryIdentifier` MUST NOT return `FAILED` with `NotFound` error code.

- A `list` request MUST support pagination by returning a `NextToken`.

The `NextToken` returned MUST be able to be used in a subsequent `list` request to retrieve the next set of results from the service.

The `NextToken` MUST be null when all results have been returned.

- A `list` request MUST return an empty array if there are no resources found.
- A `list` handler MAY accept a set of properties conforming to the shape of the resource schema as filter criteria.

The filter should use `AND(&)` when multiple properties are passed in.

Additional requirements

The following requirements also apply to resource handlers.

Returned models must conform to the shape of the schema

A model returned in a `ProgressEvent` object MUST always conform to the shape of the resource schema. This means that each property that is returned MUST adhere to its own individual restrictions: correct data type, regex, length, etc. However, the model returned MAY NOT contain all properties defined as required in the json-schema. For example, all handlers MUST NOT return `writeOnlyProperties`, which may be required as input to Create and Update handlers.

More specifically, contract tests validate models based on json-schema [Validation Keywords](#).

- ALL Validation Keywords for the following MUST be observed:
 - Any Instance Type (Section 6.1)
 - Numeric Instances (Section 6.2)
 - Strings (Section 6.3)
 - Arrays (Section 6.4)
- All Validation Keywords for Objects (Section 6.5) MUST be observed EXCEPT for:
 - `required` (Section 6.5.3)
 - `dependencies` (Section 6.5.7)
 - `propertyNames` (Section 6.5.8)
- Contract tests will NOT validate Validation Keywords for:
 - Applying Subschemas Conditionally (Section 6.6)
 - Applying Subschemas With Boolean Logic (Section 6.7)

Update, create, and delete handlers must satisfy desired-state stabilization

Stabilization is the process of waiting for a resource to be in a particular state. Note that reaching the desired-state is mandatory for all handlers before returning `SUCCESS`.

Create and update handlers

For Create and Update handlers, desired-state stabilization is satisfied when all properties specified in the request are applied as requested. This is verified by calling the Read handler.

In many cases, the desired-state is reached immediately upon completion of a Create/Update API call. However, in some cases, multiple API calls and/or wait periods may be required in order to reach this state.

Eventual consistency in desired-state stabilization

Eventual consistency means that the result of an API command you run might not be immediately visible to all subsequent commands you run. Handling API eventual consistency is required as part of desired-state stabilization. This is because a subsequent Read call might fail with a `NotFound` error code.

AWS EC2 resources are a great example of this. For more information, see [Eventual Consistency](#) in the *Amazon Elastic Compute Cloud API Reference*.

Examples of desired-state stabilization

For a simple example of desired-state stabilization, consider the implementation of the `create` handler for the `AWS::Logs::MetricFilter` resource: immediately after the handler code completes the call to the `PutMetricFilter` method, the `AWS::Logs::MetricFilter` has achieved its desired state. You can examine the code for this resource in its open-source repository at github.com/aws-cloudformation/aws-cloudformation-resource-providers-logs.

A more complex example is the implementation of the `update` handler for the `AWS::Kinesis::Stream` resource. The update handler must make multiple API calls during an update, including `AddTagsToStream` or `RemoveTagsFromStream`, `UpdateShardCount`, `IncreaseRetentionPeriod` or `DecreaseRetentionPeriod`, and `StartStreamEncryption` or `StopStreamEncryption`. Meanwhile, each API call will set the `StreamStatus` to `UPDATING`, during which time other API calls cannot be performed or the API will throw a `ResourceInUseException`. Therefore, in order to reach the desired state, the handler will need to wait for the `StreamStatus` to become `ACTIVE` in between each API call.

Delete handlers

In most cases, the definition of 'deleted' is obvious. A `delete` API call will result in the resource being purged from the database, and the resource is no longer describable to the user.

However, in some cases, a deletion will result in the resource leaving an *audit trail*, in which the resource can still be described by service APIs, but can no longer be interacted with by the user. For example, when you delete a CloudFormation stack, it is assigned a status of `DELETE_COMPLETE`, but it can still be returned from a [DescribeStacks](#) API call. For resources like this, the desired-state for deletion is when the resource has reached a *terminal, inoperable, and irrecoverable state*. If the resource can continue to be mutated by the user through another API call, then it is not *deleted*, it is *updated*.

Note that there is no difference between desired-state stabilization and runtime-state stabilization for a delete handler. By definition, once a resource has reached the desired-state for deletion, a subsequent `read` call **MUST** return `FAILED` with a `NotFound` error code, and a subsequent `create` call with the same `primaryIdentifier` or `additionalIdentifiers` **MUST NOT** return `FAILED` with an `AlreadyExists` error code. Additional restrictions are defined in the contract above.

So in the case of a CloudFormation stack, a `read` handler **MUST** return `FAILED` with a `NotFound` error code if the stack is `DELETE_COMPLETE`, even though its audit trail can still be accessed by the `DescribeStacks` API.

Update and create handlers should satisfy runtime-state stabilization

Runtime-state stabilization is a process of waiting for the resource to be "ready" to use. Generally, runtime-state stabilization is done by continually describing the resource until it reaches a particular state, though it can take many forms.

Runtime-state stabilization can mean different things for different resources, but the following are common requirements:

- *Additional mutating API calls can be made on the resource*

Some resources cannot be modified while they are in a particular state

- *Dependent resources can consume the resource*

There may be other resources which need to consume the resource in some way, but can't until it is in a particular state.

- *Users can interact with the resource*

Customers may not be able to use the resource until it is in a particular status. This usually overlaps with the dependent resources requirement, although there could be different qualifications, depending on the resources.

Note that while desired-state stabilization is mandatory, runtime-state stabilization is optional but encouraged. Users have come to expect that once a resource is `COMPLETE`, they will be able to use it.

Examples of run-time stabilization

For a simple example of desired-state stabilization, consider the implementation of the `create` handler for the `AWS::KinesisFirehose::DeliveryStream` resource. The `create` handler invokes only a single API, `CreateDeliveryStream`, in order for the resource to reach its desired state. Immediately after this API call is made, a `read` request will return the correct desired state. However, the resource still has not reached runtime-stabilization because it cannot be used by the customer or downstream resources until the `DeliveryStreamStatus` is `ACTIVE`.

For a more complex example, consider the implementation of the `update` handler for the `AWS::Kinesis::Stream` resource once again. Once the `update` handler has made its final call, to `StartStreamEncryption` or `StopStreamEncryption` as described in [Examples of desired-state stabilization](#), the resource has reached its desired state. However, like the other API calls on the Kinesis resource, the `StreamStatus` will again be set to `UPDATING`. During this period, it has reached its desired state, and customers can even continue using the stream. But it has not yet achieved runtime-stabilization, because additional API calls cannot be made on the resource until the `StreamStatus` gets set to `ACTIVE`.

Handlers must not leak resources

Resource leaking refers to when a handler loses track of the existence of a resource. This happens most often in the following cases:

- A `create` handler is not idempotent. Re-invoking the handler with the same `idempotencyToken` will cause another resource to be created, and the handler is only tracking a single resource.
- A `create` handler creates the resource, but is unable to communicate an identifier for that resource back to CloudFormation. A subsequent `delete` call does not have enough information to delete the resource.
- A bug in the `delete` handler causes the resource to not actually be deleted, but the `delete` handler reports that the resource was successfully deleted.

Handler Error Codes

One of the following error codes **MUST** be returned from the handler whenever there is a [progress event](#) with an operation status of `FAILED`.

- `AccessDenied`

The customer has insufficient permissions to perform this request.

Type: Terminal

- `AlreadyExists`

The specified resource already existed prior to the execution of this handler. This error is applicable to create handlers only.

Type: Terminal

- `GeneralServiceException`

The downstream service generated an error that does not map to any other handler error code.

Type: Terminal

- `InternalFailure`

An unexpected error occurred within the handler.

Type: Terminal

- `InvalidCredentials`

The credentials provided by the user are invalid.

Type: Terminal

- `InvalidRequest`

Invalid input from the user has generated a generic exception.

Type: Terminal

- `NetworkFailure`

The request could not be completed due to networking issues, such as a failure to receive a response from the server.

Type: Retriable

- `NotFound`

The specified resource does not exist, or is in a terminal, inoperable, and irrecoverable state.

Type: Terminal

- `NotStabilized`

The downstream resource failed to complete all of its ready-state checks.

Type: Retriable

- `NotUpdatable`

The user has requested an update to a property defined in the resource provider schema as a [create-only property](#). This error is applicable to update handlers only.

Type: Terminal

- `ResourceConflict`

The resource is temporarily unavailable to be acted upon. For example, if the resource is currently undergoing an operation and cannot be acted upon until that operation is finished.

Type: Retriable

- `ServiceInternalError`

The downstream service returned an internal error, typically with a 5XX HTTP Status code.

Type: Retriable

- `ServiceLimitExceeded`

A non-transient resource limit was reached on the service side.

Type: Terminal

- `Throttling`

The request was throttled by a downstream service. Retriable.

Type: Retriable

ProgressEvent Object Schema

A `ProgressEvent` is a JSON object which represents the current operation status of the handler, the current live state of the resource, and any additional resource information the handler wishes to communicate to the CloudFormation CLI. Each handler **MUST** communicate a progress event to the CloudFormation CLI under certain circumstances, and **SHOULD** communicate a progress event under others. For more information, see [Handler Communication Contract](#).

A handler **MAY** use progress events on a re-invocation to continue work from where it left off. For a detailed discussion of this, see [Progress Chaining, Stabilization and Callback Pattern](#).

Syntax

Below is the syntax for the `ProgressEvent` object.

```
{
  "OperationStatus": "string",
  "HandlerErrorCode": "string",
  "Message": "string",
  "CallbackContext": "string",
  "CallbackDelaySeconds": "string",
  "ResourceModel": "string",
  "ResourceModels": [
    "string"
  ],
  "NextToken": "string",
}
```

Properties

`OperationStatus`

Indicates whether the handler has reached a terminal state or is still computing and requires more time to complete.

Values: `PENDING` | `IN_PROGRESS` | `SUCCESS` | `FAILED`

Required: No

`HandlerErrorCode`

A handler error code should be provided when the event operation status is `FAILED` or `IN_PROGRESS`.

For a list of handler error codes, see [Handler Error Codes](#).

Required: Conditional. A handler error codes **MUST** be returned from the handler whenever there is a progress event with an operation status of `FAILED`.

Message

Information which can be shown to users to indicate the nature of a progress transition or callback delay.

Required: No

CallbackContext

Arbitrary information which the handler can return in an event with operation status of `IN_PROGRESS`, to allow the passing through of additional state or metadata between subsequent retries. For example, to pass through a resource identifier which can be used to continue polling for stabilization.

For more detailed examples, see [Progress Chaining, Stabilization and Callback Pattern](#).

Required: No

CallbackDelaySeconds

A callback will be scheduled with an initial delay of no less than the number of seconds specified.

Set this value to less than 0 to indicate no callback should be made.

Required: No

ResourceModel

Resource model returned by a `read` or `list` operation response for synchronous results, or for final response validation/confirmation by `create`, `update`, and `delete` operations.

Required: No

ResourceModels

List of resource models returned by a `list` operation response for synchronous results.

Required: Conditional. Required for List handlers.

NextToken

Token used to request additional pages of resources from a `list` operation response.

Required: Conditional. Required for List handlers.

Progress Chaining, Stabilization and Callback Pattern

Often when you develop CloudFormation resources, when interacting with web service APIs you need to chain them in sequence to apply the desired state. CloudFormation provides a framework to write these chain patterns. The framework does a lot of the heavy lifting needed to handle error conditions, throttle when calling downstream API, and more. The framework provides callbacks that the handler can use to inspect and change the behavior when making these service calls.

Most web service API calls follows a typical pattern:

1. Initiate the call context for the API.
2. Transform the incoming resource model properties to the underlying service API request.
3. Make the service call.
4. Handle errors. (Optional)
5. Handle stabilization. (Optional, if you need resource to be in a specific state before you apply the next state.)

6. Finalize progress to the next part of the call chain, or indicate successful completion.

In writing the handler, you do not need to do anything special with replay/continuation semantics. The framework ensures that the call chain is effectively resumed from where it was halted. This is essentially useful when the wait time for resource stabilization runs into minutes or even hours.

Sample: Kinesis Stream Integration

Here is an example integration against AWS service APIs for a Kinesis Stream. A snippet of the Kinesis resource model is shown below:

```
{
  "typeName": "AWS::Kinesis::Stream",
  "description": "Resource Type definition for AWS::Kinesis::Stream",
  "definitions": {
    ...
  },
  "properties": {
    "Arn": {
      "type": "string"
    },
    "Name": {
      "type": "string",
      "pattern": "[a-zA-Z0-9_.-]+"
    },
    "RetentionPeriodHours": {
      "type": "integer",
      "minimum": 24,
      "maximum": 168
    },
    "ShardCount": {
      "type": "integer",
      "minimum": 1,
      "maximum": 100000
    },
    "StreamEncryption": {
      "$ref": "#/definitions/AWSKinesisStreamStreamEncryption"
    },
    "Tags": {
      "type": "array",
      "uniqueItems": true,
      "items": {
        "$ref": "#/definitions/Tag"
      },
      "maximum": 50
    }
  }
  ...
}
```

And a sample CloudFormation template for creating this resource in a stack:

```
---
AWSTemplateFormatVersion: '2010-09-09'
Description: AWS MetricFilter
Resources:
  KinesisStream:
    Type: AWS::Kinesis::Stream
    Properties:
      ShardCount: 100
      RetentionPeriodHours: 36
    Tags:
```

```
- Key: '1'
  Value: one
- Key: '2'
  Value: two
StreamEncryption:
  EncryptionType: KMS
  KeyId: alias/KinesisEncryption
```

For Kinesis, the stream must first be created with a name and shard count, then tags can be applied, followed by encryption. After creating a stream, but before any other configuration can be applied, the stream must be in an ACTIVE state.

Here is the example of the using the progress-chaining and callback pattern to apply state consistently. Note that much of the error handling is delegated to the framework. The CloudFormation CLI provides some error handling on interpreting errors that can be retried after a delay. The framework provides a fluent API that guides the developer with the right set of calls with strong typing and code completion capabilities in IDEs.

```
public class CreateHandler extends BaseKinesisHandler {
    //
    // The handler is provide with a AmazonWebServicesClientProxy that provides
    // the framework for making calls that returns a ProgressEvent,
    // which can then be chained to perform the next task.
    //
    protected ProgressEvent<ResourceModel, CallbackContext>
        handleRequest(final AmazonWebServicesClientProxy proxy,
                      final ResourceHandlerRequest<ResourceModel> request,
                      final CallbackContext callbackContext,
                      final ProxyClient<KinesisClient> client,
                      final Logger logger) {

        ResourceModel model = request.getDesiredResourceState();
        if (model.getName() == null || model.getName().isEmpty()) {
            model.setName(
                IdentifierUtils.generateResourceIdentifier(
                    "stream-", request.getClientRequestToken(), 128));
        }
        //
        // 1) initiate the call context, we are making createStream API call
        //
        return proxy.initiate(
            "kinesis:CreateStream", client, model, callbackContext)

            //
            // 2) transform Resource model properties to CreateStreamRequest API
            //
            .request((m) ->
                CreateStreamRequest.builder()
                    .streamName(m.getName()).shardCount(m.getShardCount()).build())

            //
            // 3) Make a service call. Handler does not worry about credentials, they
            //     are auto injected
            //
            .call((r, c) ->
                c.injectCredentialsAndInvokeV2(r, c.client()::createStream))

            //
            // provide stabilization callback. The callback is provided with
            // the following parameters
            //   a. CreateStreamRequest the we transformed in request()
            //   b. CreateStreamResponse that the service returned with successful call
            //   c. ProxyClient<Kinesis>, we provided in initiate call
            //   d. ResourceModel we provided in initiate call
        }
```

```
// f. CallbackContext callback context.
//
//
.stabilize((_request, _response, _client, _model, _context) ->
    isStreamActive(client1, _model, context))

//
// Once ACTIVE return progress
//
.progress()

//
// we then chain to next state, setting tags on the resource.
// we receive ProgressEvent object from .progress().
//
.then(r -> {
    Set<Tag> tags = model.getTags();
    if (tags != null && !tags.isEmpty()) {
        return setTags(proxy, client, model, callbackContext, false, logger);
    }
    return r;
})

//
// we then setRetention...
//
.then(r -> {
    Integer retention = model.getRetentionPeriodHours();
    if (retention != null) {
        return handleRetention(proxy, client, model, DEFAULT_RETENTION,
retention, callbackContext, logger);
    }
    return r;
})

... // other steps

//
// finally we wait for Kinesis stream to be ACTIVE
//
.then((r) -> waitForActive(proxy, client, model, callbackContext))

//
// we then delete to ReadHandler to read the live state and send
// back successful response.
//
.then((r) -> new ReadHandler()
    .handleRequest(proxy, request, callbackContext, client, logger));
}
}
```

How to Make Other Calls

The same pattern shown here for `CreateStreamRequest` is followed with others as well. Here is code for `handleRetention`:

```
protected ProgressEvent<ResourceModel, CallbackContext>
    handleRetention(final AmazonWebServicesClientProxy proxy,
        final ProxyClient<KinesisClient> client,
        final ResourceModel model,
        final int previous,
        final int current,
        final CallbackContext callbackContext,
        final Logger logger) {
```

```
        if (current > previous) {
            //
            // 1) initiate the call context, we are making IncreaseRetentionPeriod API
call
            //
            return proxy.initiate(
                "kinesis:IncreaseRetentionPeriod:" + getClass().getSimpleName(),
                client, model, callbackContext)
            //
            // 2) transform Resource model properties to
IncreaseStreamRetentionPeriodRequest API
            //
            .request((m) ->
                IncreaseStreamRetentionPeriodRequest.builder()
                    .retentionPeriodHours(current)
                    .streamName(m.getName()).build())
            //
            // 3) Make a service call. Handler does not worry about credentials, they
            // are auto injected

            // Add important comments like shown below
            // https://docs.aws.amazon.com/kinesis/latest/APIReference/
API_IncreaseStreamRetentionPeriod.html
            // When applying change if stream is not ACTIVE, we get ResourceInUse.
            // We filter this expectation back off and then re-try to set this.

            //
            // set new retention period
            //
            .call((r, c) -> c.injectCredentialsAndInvokeV2(r,
c.client()::increaseStreamRetentionPeriod))

            //
            // Filter ResourceInUse or LimitExceeded.
            // Currently LimitExceeded is issued even for throttles
            //
            .exceptFilter(this::filterException).progress();
        } else {
            return proxy.initiate("kinesis:DecreaseRetentionPeriod:" +
getClass().getSimpleName(), client, model, callbackContext)
            //
            // convert to API model
            //
            .request(m ->
DecreaseStreamRetentionPeriodRequest.builder().retentionPeriodHours(current).streamName(m.getName())
                .build())
            ... // snipped for brevity
            .exceptFilter(this::filterException).progress();
        }
    }

protected boolean filterException(AwsRequest request,
                                Exception e,
                                ProxyClient<KinesisClient> client,
                                ResourceModel model,
                                CallbackContext context) {
    return e instanceof ResourceInUseException ||
        e instanceof LimitExceededException;
}
```

CloudFormation CLI Command Reference

The following commands are available through the CloudFormation Command Line Interface (CLI).

Global Parameters

[init](#)

[generate](#)

[validate](#)

[invoke](#)

[test](#)

[submit](#)

Global Parameters

The following parameters can be used with any CloudFormation CLI command.

`-h, --help`

Show the help message and exit.

`-v, --verbose`

Increase the output verbosity. Can be specified multiple times.

init

Description

Generates a new resource project with stub source files. While the specific folder structure and files generated varies by language, in general the project includes:

- Resource schema file
- Handler function source files
- Unit test files
- IDE and build files for the specified language

By default, `init` generates the resource project in the current directory.

Synopsis

```
cfn init  
[--force]
```

Options

`--force`

Force project files to be overwritten.

Output

The `init` command launches a wizard that walks you through setting up the project, including specifying the resource name.

generate

Description

Generates code based on the project and resource provider schema.

Synopsis

```
cfn generate
```

Output

```
Generated files for <resource_name>
```

validate

Description

Validates a project's resource specification against the [Resource Provider Definition Schema](#).

Synopsis

```
cfn validate
```

Output

invoke

Description

Performs contract tests on the specified handler of a resource provider.

Synopsis

```
cfn invoke  
[--endpoint <value>]  
[--function-name <value>]  
[--region <value>]  
[--max-reinvoke <value>]  
action  
request
```

Options

`--endpoint <value>`

The endpoint at which the type can be invoked. Alternately, you can also specify an actual Lambda endpoint and function name in your AWS account.

Default: `http://127.0.0.1:3001`

`--function-name <value>`

The logical lambda function name in the SAM template. Alternately, you can also specify an actual Lambda endpoint and function name in your AWS account.

Default: `TestEntrypoint`

`--region <value>`

The region to configure the client to interact with.

Default: `us-east-1`

`--max-reinvoke <value>`

Maximum number of IN_PROGRESS re-inocations allowed before exiting. If not specified, will continue to re- invoke until terminal status is reached.

`action`

Which single handler to invoke.

Values: `CREATE | READ | UPDATE | DELETE | LIST`

`request`

File path to a JSON file containing the request with which to invoke the function.

Output

test

Description

Performs contract tests on the handlers of a resource provider.

Synopsis

```
cfn test
[--endpoint <value>]
[--function-name <value>]
[--region <value>]
[--role-arn <value>]
```

Options

`--endpoint <value>`

The endpoint at which the type can be invoked. Alternately, you can also specify an actual Lambda endpoint and function name in your AWS account.

Default: `http://127.0.0.1.3001`

`--function-name <value>`

The logical lambda function name in the SAM template. Alternately, you can also specify an actual Lambda endpoint and function name in your AWS account.

Default: `TestEntrypoint`

`--region <value>`

The region to use for temporary credentials.

Default: `us-east-1`

`--role-arn <value>`

The Amazon Resource Name (ARN) of the IAM execution role for the contract tests to assume and use when performing operations.

If you do not specify an execution role, the contract tests use the environment credentials or the credentials specified in the [Boto 3 credentials chain](#).

Output

submit

Description

Registers the resource provider with CloudFormation, in the specified region. Registering a resource provider makes it available for use in CloudFormation operations. Registering includes:

- Validating the resource schema.
- Packaging up the resource project files and uploading them to CloudFormation.

This includes the source code for your resource handlers. These resource handlers run within the CloudFormation account.

- Determining which handlers have been specified for the resource, and running the appropriate contract tests.
- Uploading the resource handlers as functions that CloudFormation calls at the appropriate times in a resource's lifecycle.
- Returning a *registration token* that you can use with the [DescribeTypeRegistration](#) action to track the status of the registration request.

Note

As part of registering a resource provider type, CloudFormation must be able to access the S3 bucket which contains the schema handler package for that resource provider. For more information, see [IAM Permissions for Registering a Resource Provider](#) in the *AWS CloudFormation User Guide*.

Synopsis

```
cfn submit
[--dry-run]
[--endpoint-url <value>]
[--region <value>]
[--role-arn <value>]
[--no-role]
[--set-default]
```

Options

`--dry-run`

Validate the schema and package up the project files, but do not register the resource provider with CloudFormation.

`--endpoint-url <value>`

The CloudFormation endpoint to use.

`--region <value>`

The AWS region in which to register the resource provider. If no region is specified, the resource provider is registered in the default region.

`--role-arn <value>`

A specific IAM role to use when invoking handler operations.

If you do not specify an IAM role, the CloudFormation CLI attempts to create or update an execution role based on the execution role template derived from the resource type's schema, and then passes this execution role to CloudFormation. For more information, see [Accessing AWS APIs from a Resource Type](#).

You cannot specify both `--role-arn` and `--no-role` arguments.

`--no-role`

Prevent the CloudFormation CLI from passing an execution role to CloudFormation.

If your resource type calls AWS APIs in any of its handlers, you must either specify a role arn, or have the CloudFormation CLI create or update an execution role and pass that execution role to CloudFormation. For more information, see [Accessing AWS APIs from a Resource Type](#).

You cannot specify both `--role-arn` and `--no-role` arguments.

`--set-default`

Upon successful registration of the type version, sets the current type version as the default version.

Output

Resource provider registration is an asynchronous operation. You can use the supplied registration token to track the progress of your provider registration request using the [DescribeTypeRegistration](#) action of the CloudFormation API.

Registering Resource Providers for Use in AWS CloudFormation Templates

Once you've completed developing your resource provider, you'll need to *register* it with CloudFormation in order to make it available for use in CloudFormation operations. From the CloudFormation CLI, use the `submit` (p. 33) command to register your resource with CloudFormation. The `submit` command does the following:

- Validates the resource schema.
- Packages up the resource project files and uploads them to CloudFormation.

This includes the source code for your resource handlers. These resource handlers run within the CloudFormation service account.

- Runs the unit and contract tests defined in the resource project.
- Determines which handlers have been specified for the resource, to determine how CloudFormation provisions the resource.

- Returns a *registration token* that you can use with the [DescribeTypeRegistration](#) action to track the status of the registration request.

To validate and package your resource project, but not register it with CloudFormation, use the `--dry-run` option for the `submit` command.

You can also register your resource directly using the [RegisterType](#) action.

You must register a resource in each region in which you want to use it.

Use the [ListTypes](#) action for summary information about types that have been registered with CloudFormation, and the [DescribeType](#) action for detailed information about specific registered resource provider or resource provider version.

Resource Provider Versions

When you register a resource provider, you are actually registering a specific *version* of that resource provider. You can register multiple versions of a resource provider, and specify which version you want to use.

Use the [SetTypeDefaultVersion](#) action to specify the default version of a type. The default version of a resource provider will be used in CloudFormation operations.

Resource Provider Scope

Any resource provider you register is only visible and usable within the account(s) in which you register it.

Resource Provider Provisioning

During registration, CloudFormation examines which resource handlers have been implemented for the resource. The handlers implemented determine what provisioning actions CloudFormation takes with respect to the resource during various stack operations.

- If the resource provider does not contain `create`, `read`, and `delete` handlers, CloudFormation cannot actually provision the resource.
- If the resource provider does not contain an `update` handler, CloudFormation cannot update the resource during stack update operations, and will instead replace it.

Deregistering Resource Providers and Provider Versions

To remove a resource provider or provider version from active use in CloudFormation, you must *deregister* it using the [DeregisterType](#) action. If a type or type version is deregistered, it can no longer be used in CloudFormation operations.

You can deregister a specific resource provider version, or the resource provider as a whole. To deregister a type, you must individually deregister all registered versions of that type. If a type has only a single registered version, deregistering that version results in the type itself being deregistered. You cannot deregister the default version of a type, unless it is the only registered version of that type, in which case the type itself is deregistered as well.

Deregistering a resource provider or resource provider version deregisters it in all regions.

You cannot deregister a resource using the CloudFormation console.

Walkthrough: Develop a Resource Provider

In this walkthrough, we'll use the CloudFormation CLI to create a sample resource provider, `Example::Testing::WordPress`. This includes modeling the schema, developing the handlers to testing those handlers, all the way to performing a dry run to get the resource provider ready to submit to the CloudFormation registry. We'll be coding our new resource provider in Java, and using the `us-west-2` region.

Prerequisites

For purposes of this walkthrough, it is assumed you have already set up the CloudFormation CLI and associated tooling for your Java development environment:

[Setting Up Your Environment for Developing Resource Providers \(p. 1\)](#)

Create the Resource Provider Development Project

Before we can actually design and implement our resource provider, we'll need to generate a new resource type project, and then import it into our IDE.

Note

This walkthrough uses the Community Edition of the [IntelliJ IDEA](#).

Initiate the project

1. Use the `init` command to create your resource provider project and generate the files it requires.

```
$ cfn init
Initializing new project
```

2. The `init` command launches a wizard that walks you through setting up the project, including specifying the resource name. For this walkthrough, specify `Example::Testing::WordPress`.

```
Enter resource type identifier (Organization::Service::Resource):
Example::Testing::WordPress
```

The wizard then enables you to select the appropriate language plugin. Currently, the only language plugin available is for Java:

```
One language plugin found, defaulting to java
```

3. Finally, specify the package name. For this walkthrough, use `com.example.testing.wordpress`

```
Enter a package name (empty for default 'com.example.testing.wordpress'):
com.example.testing.wordpress
Initialized a new project in /workplace/tobflem/example-testing-wordpress
```

Initiating the project includes generating the files needed to develop the resource provider. For example:

```
$ ls -l
README.md
example-testing-wordpress.json
lombok.config
pom.xml
rpdk.log
```

```
src
target
template.yml
```

Import the project into your IDE

In order to guarantee that any project dependencies are correctly resolved, you must import the generated project into your IDE with Maven support.

For example, if you are using IntelliJ IDEA, you would need to do the following:

1. From the **File** menu, choose **New**, then choose **Project From Existing Sources**.
2. Navigate to the project directory
3. In the **Import Project** dialog box, choose **Import project from external model** and then choose **Maven**.
4. Choose **Next** and accept any defaults to complete importing the project.

Model the Resource Provider

When you initiate the resource provider project, an example resource provider schema file is included to help start you modeling your resource provider. This is a JSON file named for your resource, and contains an example of a typical resource provider schema. In the case of our example resource, the schema file is named `example-testing-wordpress.json`.

1. In your IDE, open `example-testing-wordpress.json`.
2. Paste the following schema in place of the default example schema currently in the file.

This schema defines a resource, `Example::Testing::WordPress`, that provisions a WordPress site. The resource itself contains four properties, only two of which can be set by users: `Name`, and `SubnetId`. The other two properties, `InstanceId` and `PublicIp`, are read-only, meaning they cannot be set by users, but will be assigned during resource creation. Both of these properties also serve as identifiers for the resource when it is provisioned.

As we'll see later in the walkthrough, creating a WordPress site actually requires more information than represented in our resource model. However, we'll be handling that information on behalf of the user in the code for the resource `create` handler.

```
{
  "typeName": "Example::Testing::WordPress",
  "description": "An example resource that creates a website based on WordPress 5.2.2.",
  "sourceUrl": "https://docs.aws.amazon.com/cloudformation-cli/latest/userguide/resource-
type-walkthrough.html",
  "properties": {
    "Name": {
      "description": "A name associated with the website.",
      "type": "string",
      "pattern": "^[a-zA-Z0-9]{1,219}\\Z",
      "minLength": 1,
      "maxLength": 219
    },
    "SubnetId": {
      "description": "A subnet in which to host the website.",
      "pattern": "^(subnet-[a-f0-9]{13})|(subnet-[a-f0-9]{8})\\Z",
      "type": "string"
    },
    "InstanceId": {
      "description": "The ID of the instance that backs the WordPress site.",
      "type": "string"
    }
  }
}
```



```
    },
    "PublicIp": {
      "description": "The public IP for the WordPress site.",
      "type": "string"
    }
  },
  "required": [
    "Name",
    "SubnetId"
  ],
  "handlers": {
    "create": {
      "permissions": [
        "ec2:AuthorizeSecurityGroupIngress",
        "ec2:CreateSecurityGroup",
        "ec2>DeleteSecurityGroup",
        "ec2:DescribeInstances",
        "ec2:DescribeSubnets",
        "ec2:CreateTags",
        "ec2:RunInstances"
      ]
    },
    "read": {
      "permissions": [
        "ec2:DescribeInstances"
      ]
    },
    "delete": {
      "permissions": [
        "ec2>DeleteSecurityGroup",
        "ec2:DescribeInstances",
        "ec2:TerminateInstances"
      ]
    }
  },
  "additionalProperties": false,
  "primaryIdentifier": [
    "/properties/PublicIp",
    "/properties/InstanceId"
  ],
  "readOnlyProperties": [
    "/properties/PublicIp",
    "/properties/InstanceId"
  ]
}
```

3. Update the auto-generated files in the resource provider package so that they reflect the changes we've made to the resource provider schema.

When we first initiated the resource provider project, the CloudFormation CLI generated supporting files and code for our resource provider. Since we've made changes to the resource provider schema, we'll need to regenerate that code to ensure that it reflects the updated schema. To do this, we use the generate command:

```
$ cfn generate
Generated files for Example::Testing::WordPress
```

Note

When using Maven, as part of the build process the generate command is automatically run before the code is compiled. So your changes will never get out of sync with the generated code.

Be aware the CloudFormation CLI must be in a location Maven/the system can find. For more information, see [Setting Up Your Environment for Developing Resource Providers \(p. 1\)](#).

Implement the Resource Handlers

Now that we have our resource provider schema specified, we can start implementing the behavior we want the resource provider to exhibit during each resource operation. To do this, we'll have to implement the various event handlers, including:

- Adding any necessary dependencies
- Writing code to implement the various resource operation handlers.

Add Dependencies

To actually make WordPress handlers that call the associated EC2 APIs, we need to declare the AWS EC2 SDK as a dependency in Maven's `pom.xml` file. To enable this, we need to add a dependency on the [AWS SDK for Java](#) to the project.

1. In your IDE, open the project's `pom.xml` file.
2. Add the following dependency in the `dependencies` section.

```
<dependency>
  <groupId>com.amazonaws</groupId>
  <artifactId>aws-java-sdk-ec2</artifactId>
  <version>1.11.606</version>
</dependency>
```

This artifact will be added by Maven from the [Maven Repository](#).

For more information on how to add dependencies, see the [Maven documentation](#).

Note

Depending on your IDE, you may have to take additional steps for your IDE to include the new dependency.

In IntelliJ IDEA, a dialog should appear to enable you to import these changes. we recommend allowing automatic importing.

Implement the Create Handler

With the necessary dependency specified, we can now start writing the handlers that actually implement the resource's functionality. For our example resource, we'll implement just the `create` and `delete` operation handlers.

To create a WordPress site, our resource `create` handler will have to accomplish the following:

- Gather and define inputs that we'll need to create the underlying AWS resources on behalf of the user. These are details we're managing for them, since this is a very high-level resource type.
- Create an EC2 instance using a special AMI vended by Bitnami from the AMI Marketplace that bootstraps WordPress.
- Create a security group that the instance will belong to so you can access the WordPress site from your browser.
- Change the security group rules to dictate what the networking rules are for web access to the WordPress site.
- If something goes wrong with creating the resource, attempt to delete the security group.

Define the CallbackContext

Because our create handler is more complex than simply calling a single API, it takes some time to complete. However, each handler times out after one minute. To work around this issue, we'll write our handlers as state machines. A handler can exit with one of three states: SUCCESS, IN_PROGRESS, and FAILED. To wait on stabilization of underlying resources, we can return an IN_PROGRESS state with a CallbackContext. The CallbackContext will hold details about the current state of the execution. When we return an IN_PROGRESS state and a CallbackContext, CloudFormation will re-invoke the handler and pass the CallbackContext in with the request. You can then make decisions based on what is included in the context.

The CallbackContext is modeled as a POJO so you can define what information you want to pass between state transitions explicitly.

1. In your IDE, open the `CallbackContext.java` file, located in the `src/main/java/com/example/testing/wordpress` folder.
2. Replace the entire contents of the `CallbackContext.java` file with the following code.

```
package com.example.testing.wordpress;

import com.amazonaws.services.ec2.model.Instance;

import java.util.List;

import lombok.AllArgsConstructor;
import lombok.Builder;
import lombok.Data;
import lombok.NoArgsConstructor;
import lombok.NoArgsConstructor;

@Builder(toBuilder = true)
@Data
@NoArgsConstructor
@AllArgsConstructor
public class CallbackContext {
    private Instance instance;
    private Integer stabilizationRetriesRemaining;
    private List<String> instanceSecurityGroups;
}
```

Code the Create Handler

1. In your IDE, open the `CreateHandler.java` file, located in the `src/main/java/com/example/testing/wordpress/CreateHandler.java` folder.
2. Replace the entire contents of the `CreateHandler.java` file with the following code.

```
package com.example.testing.wordpress;

import software.amazon.cloudformation.proxy.AmazonWebServicesClientProxy;
import software.amazon.cloudformation.proxy.Logger;
import software.amazon.cloudformation.proxy.OperationStatus;
import software.amazon.cloudformation.proxy.ProgressEvent;
import software.amazon.cloudformation.proxy.ResourceHandlerRequest;
import com.amazonaws.services.ec2.AmazonEC2;
import com.amazonaws.services.ec2.AmazonEC2ClientBuilder;
import com.amazonaws.services.ec2.model.AuthorizeSecurityGroupIngressRequest;
import com.amazonaws.services.ec2.model.CreateSecurityGroupRequest;
import com.amazonaws.services.ec2.model.DeleteSecurityGroupRequest;
import com.amazonaws.services.ec2.model.DescribeInstancesRequest;
import com.amazonaws.services.ec2.model.DescribeInstancesResult;
```

```
import com.amazonaws.services.ec2.model.DescribeSubnetsRequest;
import com.amazonaws.services.ec2.model.DescribeSubnetsResult;
import com.amazonaws.services.ec2.model.Instance;
import com.amazonaws.services.ec2.model.InstanceNetworkInterfaceSpecification;
import com.amazonaws.services.ec2.model.IpPermission;
import com.amazonaws.services.ec2.model.IpRange;
import com.amazonaws.services.ec2.model.Reservation;
import com.amazonaws.services.ec2.model.RunInstancesRequest;
import com.amazonaws.services.ec2.model.Subnet;
import com.amazonaws.services.ec2.model.Tag;
import com.amazonaws.services.ec2.model.TagSpecification;

import java.util.List;
import java.util.UUID;

public class CreateHandler extends BaseHandler<CallbackContext> {
    private static final String SUPPORTED_REGION = "us-west-2";
    private static final String WORDPRESS_AMI_ID = "ami-04fb0368671b6f138";
    private static final String INSTANCE_TYPE = "m4.large";
    private static final String SITE_NAME_TAG_KEY = "Name";
    private static final String AVAILABLE_INSTANCE_STATE = "running";
    private static final int NUMBER_OF_STATE_POLL_RETRIES = 60;
    private static final int POLL_RETRY_DELAY_IN_MS = 5000;
    private static final String TIMED_OUT_MESSAGE = "Timed out waiting for instance to
become available.";

    private AmazonWebServicesClientProxy clientProxy;
    private AmazonEC2 ec2Client;

    @Override
    public ProgressEvent<ResourceModel, CallbackContext> handleRequest(
        final AmazonWebServicesClientProxy proxy,
        final ResourceHandlerRequest<ResourceModel> request,
        final CallbackContext callbackContext,
        final Logger logger) {

        final ResourceModel model = request.getDesiredResourceState();

        clientProxy = proxy;
        ec2Client =
AmazonEC2ClientBuilder.standard().withRegion(SUPPORTED_REGION).build();
        final CallbackContext currentContext = callbackContext == null ?

CallbackContext.builder().stabilizationRetriesRemaining(NUMBER_OF_STATE_POLL_RETRIES).build() :
callbackContext;

        // This Lambda will continually be re-invoked with the current state of the
instance, finally succeeding when state stabilizes.
        return createInstanceAndUpdateProgress(model, currentContext);
    }

    private ProgressEvent<ResourceModel, CallbackContext>
createInstanceAndUpdateProgress(ResourceModel model, CallbackContext callbackContext) {
        // This Lambda will continually be re-invoked with the current state of the
instance, finally succeeding when state stabilizes.
        final Instance instanceStateSoFar = callbackContext.getInstance();

        if (callbackContext.getStabilizationRetriesRemaining() == 0) {
            throw new RuntimeException(TIMED_OUT_MESSAGE);
        }

        if (instanceStateSoFar == null) {
            return ProgressEvent.<ResourceModel, CallbackContext>builder()
                .resourceModel(model)
                .status(OperationStatus.IN_PROGRESS)
                .callbackContext(CallbackContext.builder()
```

```

        .instance(createEC2Instance(model))

    .stabilizationRetriesRemaining(NUMBER_OF_STATE_POLL_RETRIES)
        .build()

    } else if
    (instanceStateSoFar.getState().getName().equals(AVAILABLE_INSTANCE_STATE)) {
        model.setInstanceId(instanceStateSoFar.getInstanceId());
        model.setPublicIp(instanceStateSoFar.getPublicIpAddress());
        return ProgressEvent.<ResourceModel, CallbackContext>builder()
            .resourceModel(model)
            .status(OperationStatus.SUCCESS)
            .build();

    } else {
        try {
            Thread.sleep(POLL_RETRY_DELAY_IN_MS);
        } catch (InterruptedException e) {
            throw new RuntimeException(e);
        }
        return ProgressEvent.<ResourceModel, CallbackContext>builder()
            .resourceModel(model)
            .status(OperationStatus.IN_PROGRESS)
            .callbackContext(CallbackContext.builder())

    .instance(updatedInstanceProgress(instanceStateSoFar.getInstanceId()))

    .stabilizationRetriesRemaining(callbackContext.getStabilizationRetriesRemaining() - 1)
        .build()

    }

}

private Instance createEC2Instance(ResourceModel model) {
    final String securityGroupId = createSecurityGroupForInstance(model);
    final RunInstancesRequest runInstancesRequest = new RunInstancesRequest()
        .withInstanceType(INSTANCE_TYPE)
        .withImageId(WORDPRESS_AMI_ID)
        .withNetworkInterfaces(new InstanceNetworkInterfaceSpecification()
            .withAssociatePublicIpAddress(true)
            .withDeviceIndex(0)
            .withGroups(securityGroupId)
            .withSubnetId(model.getSubnetId()))

        .withMaxCount(1)
        .withMinCount(1)
        .withTagSpecifications(buildTagFromSiteName(model.getName()));

    try {
        return clientProxy.injectCredentialsAndInvoke(runInstancesRequest,
ec2Client::runInstances)
            .getReservation()
            .getInstances()
            .stream()
            .findFirst()
            .orElse(new Instance());
    } catch (Throwable e) {
        attemptToCleanUpSecurityGroup(securityGroupId);
        throw new RuntimeException(e);
    }
}

private String createSecurityGroupForInstance(ResourceModel model) {
    String vpcId;
    try {
        vpcId = getVpcIdFromSubnetId(model.getSubnetId());
    } catch (Throwable e) {

```

```
        throw new RuntimeException(e);
    }

    final String securityGroupName = model.getName() + "-" +
UUID.randomUUID().toString();

    final CreateSecurityGroupRequest createSecurityGroupRequest = new
CreateSecurityGroupRequest()
        .withGroupName(securityGroupName)
        .withDescription("Created for the test WordPress blog: " + model.getName())
        .withVpcId(vpcId);

    final String securityGroupId =
        clientProxy.injectCredentialsAndInvoke(createSecurityGroupRequest,
ec2Client::createSecurityGroup)
            .getGroupId();

    final AuthorizeSecurityGroupIngressRequest authorizeSecurityGroupIngressRequest =
new AuthorizeSecurityGroupIngressRequest()
        .withGroupId(securityGroupId)
        .withIpPermissions(openHTTP(), openHTTPS());

    clientProxy.injectCredentialsAndInvoke(authorizeSecurityGroupIngressRequest,
ec2Client::authorizeSecurityGroupIngress);

    return securityGroupId;
}

private String getVpcIdFromSubnetId(String subnetId) throws Throwable {
    final DescribeSubnetsRequest describeSubnetsRequest = new
DescribeSubnetsRequest()
        .withSubnetIds(subnetId);

    final DescribeSubnetsResult describeSubnetsResult =
        clientProxy.injectCredentialsAndInvoke(describeSubnetsRequest,
ec2Client::describeSubnets);

    return describeSubnetsResult.getSubnets()
        .stream()
        .map(Subnet::getVpcId)
        .findFirst()
        .orElseThrow(() -> {
            throw new RuntimeException("Subnet " + subnetId +
" not found");
        });
}

private IpPermission openHTTP() {
    return new IpPermission().withIpProtocol("tcp")
        .withFromPort(80)
        .withToPort(80)
        .withIpv4Ranges(new IpRange().withCidrIp("0.0.0.0/0"));
}

private IpPermission openHTTPS() {
    return new IpPermission().withIpProtocol("tcp")
        .withFromPort(443)
        .withToPort(443)
        .withIpv4Ranges(new IpRange().withCidrIp("0.0.0.0/0"));
}

private TagSpecification buildTagFromSiteName(String siteName) {
    return new TagSpecification()
        .withResourceType("instance")
        .withTags(new Tag().withKey(SITE_NAME_TAG_KEY).withValue(siteName));
}
```

```
private Instance updatedInstanceProgress(String instanceId) {
    DescribeInstancesRequest describeInstancesRequest;
    DescribeInstancesResult describeInstancesResult;

    describeInstancesRequest = new
DescribeInstancesRequest().withInstanceIds(instanceId);
    describeInstancesResult =
clientProxy.injectCredentialsAndInvoke(describeInstancesRequest,
ec2Client::describeInstances);
    return describeInstancesResult.getReservations()
                                .stream()
                                .map(Reservation::getInstances)
                                .flatMap(List::stream)
                                .findFirst()
                                .orElse(new Instance());
}

private void attemptToCleanUpSecurityGroup(String securityGroupId) {
    final DeleteSecurityGroupRequest deleteSecurityGroupRequest = new
DeleteSecurityGroupRequest().withGroupId(securityGroupId);
    clientProxy.injectCredentialsAndInvoke(deleteSecurityGroupRequest,
ec2Client::deleteSecurityGroup);
}
}
```

Update the Create Handler Unit Test

Since our resource type is a high-level abstraction, a lot of implementation behavior isn't apparent by the name alone. As such, we'll need to make some additions to our unit tests so that we're not calling the live APIs that are necessary to create the WordPress site.

1. In your IDE, open the `CreateHandlerTest.java` file, located in the `src/test/java/com/example/testing/wordpress` folder.
2. Replace the entire contents of the `CreateHandlerTest.java` file with the following code.

```
package com.example.testing.wordpress;

import software.amazon.cloudformation.proxy.AmazonWebServicesClientProxy;
import software.amazon.cloudformation.proxy.Logger;
import software.amazon.cloudformation.proxy.OperationStatus;
import software.amazon.cloudformation.proxy.ProgressEvent;
import software.amazon.cloudformation.proxy.ResourceHandlerRequest;
import com.amazonaws.services.ec2.model.AuthorizeSecurityGroupIngressRequest;
import com.amazonaws.services.ec2.model.AuthorizeSecurityGroupIngressResult;
import com.amazonaws.services.ec2.model.CreateSecurityGroupRequest;
import com.amazonaws.services.ec2.model.CreateSecurityGroupResult;
import com.amazonaws.services.ec2.model.DescribeInstancesRequest;
import com.amazonaws.services.ec2.model.DescribeInstancesResult;
import com.amazonaws.services.ec2.model.DescribeSubnetsRequest;
import com.amazonaws.services.ec2.model.DescribeSubnetsResult;
import com.amazonaws.services.ec2.model.GroupIdentifier;
import com.amazonaws.services.ec2.model.Instance;
import com.amazonaws.services.ec2.model.InstanceState;
import com.amazonaws.services.ec2.model.Reservation;
import com.amazonaws.services.ec2.model.RunInstancesRequest;
import com.amazonaws.services.ec2.model.RunInstancesResult;
import com.amazonaws.services.ec2.model.Subnet;

import org.junit.jupiter.api.BeforeEach;
import org.junit.jupiter.api.Test;
import org.junit.jupiter.api.extension.ExtendWith;
```

```
import org.mockito.Mock;
import org.mockito.junit.jupiter.MockitoExtension;

import static org.assertj.core.api.Assertions.assertThat;
import static org.mockito.ArgumentMatchers.any;
import static org.mockito.Mockito.doReturn;
import static org.mockito.Mockito.mock;

@ExtendWith(MockitoExtension.class)
public class CreateHandlerTest {
    private String EXPECTED_TIMEOUT_MESSAGE = "Timed out waiting for instance to
    become available.";

    @Mock
    private AmazonWebServicesClientProxy proxy;

    @Mock
    private Logger logger;

    @BeforeEach
    public void setup() {
        proxy = mock(AmazonWebServicesClientProxy.class);
        logger = mock(Logger.class);
    }

    @Test
    public void testSuccessState() {
        final InstanceState inProgressState = new InstanceState().withName("running");
        final GroupIdentifier group = new GroupIdentifier().withGroupId("sg-1234");
        final Instance instance = new
Instance().withInstanceId("i-1234").withState(inProgressState).withPublicIpAddress("54.0.0.0").with

        final CreateHandler handler = new CreateHandler();

        final ResourceModel model = ResourceModel.builder()
            .name("MyWordPressSite")
            .subnetId("subnet-1234")
            .build();

        final ResourceModel desiredOutputModel = ResourceModel.builder()
            .instanceId("i-1234")
            .publicIp("54.0.0.0")
            .name("MyWordPressSite")
            .subnetId("subnet-1234")
            .build();

        final ResourceHandlerRequest<ResourceModel> request =
ResourceHandlerRequest.<ResourceModel>builder()
            .desiredResourceState(model)
            .build();

        final CallbackContext context = CallbackContext.builder()
            .stabilizationRetriesRemaining(1)
            .instance(instance)
            .build();

        final ProgressEvent<ResourceModel, CallbackContext> response
            = handler.handleRequest(proxy, request, context, logger);

        assertThat(response).isNotNull();
        assertThat(response.getStatus()).isEqualTo(OperationStatus.SUCCESS);
        assertThat(response.getCallbackContext()).isNull();
        assertThat(response.getCallbackDelaySeconds()).isEqualTo(0);
        assertThat(response.getResourceModel()).isEqualTo(desiredOutputModel);
        assertThat(response.getResourceModels()).isNull();
    }
}
```



```

        assertThat(response.getMessage()).isNull();
        assertThat(response.getErrorCode()).isNull();
    }

    @Test
    public void testInProgressStateInstanceCreationNotInvoked() {
        final InstanceState inProgressState = new InstanceState().withName("in-
progress");
        final GroupIdentifier group = new GroupIdentifier().withGroupId("sg-1234");
        final Instance instance = new
Instance().withState(inProgressState).withPublicIpAddress("54.0.0.0").withSecurityGroups(group);
        doReturn(new DescribeSubnetsResult().withSubnets(new
Subnet().withVpcId("vpc-1234")))
        .when(proxy).injectCredentialsAndInvoke(any(DescribeSubnetsRequest.class),
any());
        doReturn(new RunInstancesResult().withReservation(new
Reservation().withInstances(instance)))
        .when(proxy).injectCredentialsAndInvoke(any(RunInstancesRequest.class),
any());
        doReturn(new
CreateSecurityGroupResult().withGroupId("sg-1234"))
        .when(proxy).injectCredentialsAndInvoke(any(CreateSecurityGroupRequest.class),
any());
        doReturn(new
AuthorizeSecurityGroupIngressResult())
        .when(proxy).injectCredentialsAndInvoke(any(AuthorizeSecurityGroupIngressRequest.class),
any());

        final CreateHandler handler = new CreateHandler();

        final ResourceModel model =
ResourceModel.builder().name("MyWordPressSite").subnetId("subnet-1234").build();

        final ResourceHandlerRequest<ResourceModel> request =
ResourceHandlerRequest.<ResourceModel>builder()
            .desiredResourceState(model)
            .build();

        final ProgressEvent<ResourceModel, CallbackContext> response
            = handler.handleRequest(proxy, request, null, logger);

        final CallbackContext desiredOutputContext = CallbackContext.builder()
            .stabilizationRetriesRemaining(60)
            .instance(instance)
            .build();

        assertThat(response).isNotNull();
        assertThat(response.getStatus()).isEqualTo(OperationStatus.IN_PROGRESS);

        assertThat(response.getCallbackContext()).isEqualToComparingFieldByField(desiredOutputContext);
        assertThat(response.getCallbackDelaySeconds()).isEqualTo(0);

        assertThat(response.getResourceModel()).isEqualTo(request.getDesiredResourceState());
        assertThat(response.getResourceModels()).isNull();
        assertThat(response.getMessage()).isNull();
        assertThat(response.getErrorCode()).isNull();
    }

    @Test
    public void testInProgressStateInstanceCreationInvoked() {
        final InstanceState inProgressState = new InstanceState().withName("in-
progress");
        final GroupIdentifier group = new GroupIdentifier().withGroupId("sg-1234");
        final Instance instance = new
Instance().withState(inProgressState).withPublicIpAddress("54.0.0.0").withSecurityGroups(group);
        final DescribeInstancesResult describeInstancesResult =
            new DescribeInstancesResult().withReservations(new
Reservation().withInstances(instance));
    }

```

```
doReturn(describeInstancesResult).when(proxy).injectCredentialsAndInvoke(any(DescribeInstancesRequest.class),
any());

    final CreateHandler handler = new CreateHandler();

    final ResourceModel model =
ResourceModel.builder().name("MyWordPressSite").subnetId("subnet-1234").build();

    final ResourceHandlerRequest<ResourceModel> request =
ResourceHandlerRequest.<ResourceModel>builder()
        .desiredResourceState(model)
        .build();

    final CallbackContext context = CallbackContext.builder()
        .stabilizationRetriesRemaining(60)
        .instance(instance)
        .build();

    final ProgressEvent<ResourceModel, CallbackContext> response
        = handler.handleRequest(proxy, request, context, logger);

    final CallbackContext desiredOutputContext = CallbackContext.builder()
        .stabilizationRetriesRemaining(59)
        .instance(instance)
        .build();

    assertThat(response).isNotNull();
    assertThat(response.getStatus()).isEqualTo(OperationStatus.IN_PROGRESS);

    assertThat(response.getCallbackContext()).isEqualToComparingFieldByField(desiredOutputContext);
    assertThat(response.getCallbackDelaySeconds()).isEqualTo(0);

    assertThat(response.getResourceModel()).isEqualTo(request.getDesiredResourceState());
    assertThat(response.getResourceModels()).isNull();
    assertThat(response.getMessage()).isNull();
    assertThat(response.getErrorCode()).isNull();
}

@Test
public void testStabilizationTimeout() {
    final CreateHandler handler = new CreateHandler();

    final ResourceModel model =
ResourceModel.builder().name("MyWordPressSite").subnetId("subnet-1234").build();

    final ResourceHandlerRequest<ResourceModel> request =
ResourceHandlerRequest.<ResourceModel>builder()
        .desiredResourceState(model)
        .build();

    final CallbackContext context = CallbackContext.builder()
        .stabilizationRetriesRemaining(0)
        .instance(new
Instance().withState(new InstanceState().withName("in-progress")))
        .build();

    try {
        handler.handleRequest(proxy, request, context, logger);
    } catch (RuntimeException e) {
        assertThat(e.getMessage()).isEqualTo(EXPECTED_TIMEOUT_MESSAGE);
    }
}
}
```

Implement the Delete Handler

We'll also need to implement a delete handler. At a high level, the delete handler needs to accomplish the following:

1. Find the security groups attached to the EC2 instance that is hosting the WordPress page.
2. Delete the instance.
3. Delete the security groups.

Again, we'll implement the delete handler as a state machine.

Code the Delete Handler

1. In your IDE, open the `DeleteHandler.java` file, located in the `src/main/java/com/example/testing/wordpress` folder.
2. Replace the entire contents of the `DeleteHandler.java` file with the following code.

```
package com.example.testing.wordpress;

import software.amazon.cloudformation.proxy.AmazonWebServicesClientProxy;
import software.amazon.cloudformation.proxy.HandlerErrorCode;
import software.amazon.cloudformation.proxy.Logger;
import software.amazon.cloudformation.proxy.OperationStatus;
import software.amazon.cloudformation.proxy.ProgressEvent;
import software.amazon.cloudformation.proxy.ResourceHandlerRequest;
import com.amazonaws.services.ec2.AmazonEC2;
import com.amazonaws.services.ec2.AmazonEC2ClientBuilder;
import com.amazonaws.services.ec2.model.DeleteSecurityGroupRequest;
import com.amazonaws.services.ec2.model.DescribeInstancesRequest;
import com.amazonaws.services.ec2.model.DescribeInstancesResult;
import com.amazonaws.services.ec2.model.GroupIdentifier;
import com.amazonaws.services.ec2.model.Instance;
import com.amazonaws.services.ec2.model.Reservation;
import com.amazonaws.services.ec2.model.TerminateInstancesRequest;

import java.util.List;
import java.util.stream.Collectors;

public class DeleteHandler extends BaseHandler<CallbackContext> {
    private static final String SUPPORTED_REGION = "us-west-2";
    private static final String DELETED_INSTANCE_STATE = "terminated";
    private static final int NUMBER_OF_STATE_POLL_RETRIES = 60;
    private static final int POLL_RETRY_DELAY_IN_MS = 5000;
    private static final String TIMED_OUT_MESSAGE = "Timed out waiting for instance to terminate.";

    private AmazonWebServicesClientProxy clientProxy;
    private AmazonEC2 ec2Client;

    @Override
    public ProgressEvent<ResourceModel, CallbackContext> handleRequest (
        final AmazonWebServicesClientProxy proxy,
        final ResourceHandlerRequest<ResourceModel> request,
        final CallbackContext callbackContext,
        final Logger logger) {

        final ResourceModel model = request.getDesiredResourceState();

        clientProxy = proxy;
        ec2Client =
            AmazonEC2ClientBuilder.standard().withRegion(SUPPORTED_REGION).build();
```

```
        final CallbackContext currentContext = callbackContext == null ?
CallbackContext.builder().stabilizationRetriesRemaining(NUMBER_OF_STATE_POLL_RETRIES).build() :
                                callbackContext;

        // This Lambda will continually be re-invoked with the current state of the
        // instance, finally succeeding when state stabilizes.
        return deleteInstanceAndUpdateProgress(model, currentContext);
    }

    private ProgressEvent<ResourceModel, CallbackContext>
deleteInstanceAndUpdateProgress(ResourceModel model, CallbackContext callbackContext) {

        if (callbackContext.getStabilizationRetriesRemaining() == 0) {
            throw new RuntimeException(TIMED_OUT_MESSAGE);
        }

        if (callbackContext.getInstanceSecurityGroups() == null) {
            final Instance currentInstanceState =
currentInstanceState(model.getInstanceId());

            if (DELETED_INSTANCE_STATE.equals(currentInstanceState.getState().getName()))
{
                return ProgressEvent.<ResourceModel, CallbackContext>builder()
                    .status(OperationStatus.FAILED)
                    .errorCode(HandlerErrorCode.NotFound)
                    .build();
            }

            final List<String> instanceSecurityGroups = currentInstanceState
                .getSecurityGroups()
                .stream()
                .map(GroupIdentifier::getGroupId)
                .collect(Collectors.toList());

            return ProgressEvent.<ResourceModel, CallbackContext>builder()
                .resourceModel(model)
                .status(OperationStatus.IN_PROGRESS)
                .callbackContext(CallbackContext.builder()

.stabilizationRetriesRemaining(NUMBER_OF_STATE_POLL_RETRIES)

.instanceSecurityGroups(instanceSecurityGroups)
                                .build())
                .build();
        }

        if (callbackContext.getInstance() == null) {
            return ProgressEvent.<ResourceModel, CallbackContext>builder()
                .resourceModel(model)
                .status(OperationStatus.IN_PROGRESS)
                .callbackContext(CallbackContext.builder()

.instance(deleteInstance(model.getInstanceId()))

.instanceSecurityGroups(callbackContext.getInstanceSecurityGroups())

.stabilizationRetriesRemaining(NUMBER_OF_STATE_POLL_RETRIES)
                                .build())
                .build();
        } else if
(callbackContext.getInstance().getState().getName().equals(DELETED_INSTANCE_STATE)) {
callbackContext.getInstanceSecurityGroups().forEach(this::deleteSecurityGroup);
            return ProgressEvent.<ResourceModel, CallbackContext>builder()
                .resourceModel(model)
```

```
        .status(OperationStatus.SUCCESS)
        .build();
    } else {
        try {
            Thread.sleep(POLL_RETRY_DELAY_IN_MS);
        } catch (InterruptedException e) {
            throw new RuntimeException(e);
        }
        return ProgressEvent.<ResourceModel, CallbackContext>builder()
            .resourceModel(model)
            .status(OperationStatus.IN_PROGRESS)
            .callbackContext(CallbackContext.builder()

.instance(currentInstanceState(model.getInstanceId()))

.instanceSecurityGroups(callbackContext.getInstanceSecurityGroups())

.stabilizationRetriesRemaining(callbackContext.getStabilizationRetriesRemaining() - 1)
            .build())
        .build();
    }
}

private Instance deleteInstance(String instanceId) {
    final TerminateInstancesRequest terminateInstancesRequest = new
TerminateInstancesRequest().withInstanceIds(instanceId);
    return clientProxy.injectCredentialsAndInvoke(terminateInstancesRequest,
ec2Client::terminateInstances)
        .getTerminatingInstances()
        .stream()
        .map(instance -> new
Instance().withState(instance.getCurrentState()).withInstanceId(instance.getInstanceId()))
        .findFirst()
        .orElse(new Instance());
}

private Instance currentInstanceState(String instanceId) {
    DescribeInstancesRequest describeInstancesRequest;
    DescribeInstancesResult describeInstancesResult;

    describeInstancesRequest = new
DescribeInstancesRequest().withInstanceIds(instanceId);
    describeInstancesResult =
clientProxy.injectCredentialsAndInvoke(describeInstancesRequest,
ec2Client::describeInstances);
    return describeInstancesResult.getReservations()
        .stream()
        .map(Reservation::getInstances)
        .flatMap(List::stream)
        .findFirst()
        .orElse(new Instance());
}

private void deleteSecurityGroup(String securityGroupId) {
    final DeleteSecurityGroupRequest deleteSecurityGroupRequest = new
DeleteSecurityGroupRequest().withGroupId(securityGroupId);
    clientProxy.injectCredentialsAndInvoke(deleteSecurityGroupRequest,
ec2Client::deleteSecurityGroup);
}
}
```

Update the Delete Handler Unit Test

We'll also need to update the unit test for the delete handler.

1. In your IDE, open the `DeleteHandlerTest.java` file, located in the `src/test/java/com/example/testing/wordpress` folder.
2. Replace the entire contents of the `DeleteHandlerTest.java` file with the following code.

```
package com.example.testing.wordpress;

import software.amazon.cloudformation.proxy.AmazonWebServicesClientProxy;
import software.amazon.cloudformation.proxy.HandlerErrorCode;
import software.amazon.cloudformation.proxy.Logger;
import software.amazon.cloudformation.proxy.OperationStatus;
import software.amazon.cloudformation.proxy.ProgressEvent;
import software.amazon.cloudformation.proxy.ResourceHandlerRequest;
import com.amazonaws.services.ec2.model.DeleteSecurityGroupRequest;
import com.amazonaws.services.ec2.model.DeleteSecurityGroupResult;
import com.amazonaws.services.ec2.model.DescribeInstancesRequest;
import com.amazonaws.services.ec2.model.DescribeInstancesResult;
import com.amazonaws.services.ec2.model.GroupIdentifier;
import com.amazonaws.services.ec2.model.Instance;
import com.amazonaws.services.ec2.model.InstanceState;
import com.amazonaws.services.ec2.model.InstanceStateChange;
import com.amazonaws.services.ec2.model.Reservation;
import com.amazonaws.services.ec2.model.TerminateInstancesRequest;
import com.amazonaws.services.ec2.model.TerminateInstancesResult;

import org.junit.jupiter.api.BeforeEach;
import org.junit.jupiter.api.Test;
import org.junit.jupiter.api.extension.ExtendWith;
import org.mockito.Mock;
import org.mockito.junit.jupiter.MockitoExtension;

import java.util.Arrays;

import static org.assertj.core.api.Assertions.assertThat;
import static org.mockito.ArgumentMatchers.any;
import static org.mockito.Mockito.doReturn;
import static org.mockito.Mockito.mock;

@ExtendWith(MockitoExtension.class)
public class DeleteHandlerTest {
    private static String EXPECTED_TIMEOUT_MESSAGE = "Timed out waiting for instance to terminate.";

    @Mock
    private AmazonWebServicesClientProxy proxy;

    @Mock
    private Logger logger;

    @BeforeEach
    public void setup() {
        proxy = mock(AmazonWebServicesClientProxy.class);
        logger = mock(Logger.class);
    }

    @Test
    public void testSuccessState() {
        final DeleteSecurityGroupResult deleteSecurityGroupResult = new
DeleteSecurityGroupResult();
```

```
doReturn(deleteSecurityGroupResult).when(proxy).injectCredentialsAndInvoke(any(DeleteSecurityGroupRequest.class), any());

    final DeleteHandler handler = new DeleteHandler();

    final ResourceModel model = ResourceModel.builder().instanceId("i-1234").build();

    final ResourceHandlerRequest<ResourceModel> request =
ResourceHandlerRequest.<ResourceModel>builder()
        .desiredResourceState(model)
        .build();

    final CallbackContext context = CallbackContext.builder()
                                                .stabilizationRetriesRemaining(1)

.ininstanceSecurityGroups(Arrays.asList("sg-1234"))
Instance().withState(new InstanceState().withName("terminated")))
        .build();

    final ProgressEvent<ResourceModel, CallbackContext> response
        = handler.handleRequest(proxy, request, context, logger);

    assertThat(response).isNotNull();
    assertThat(response.getStatus()).isEqualTo(OperationStatus.SUCCESS);
    assertThat(response.getCallbackContext()).isNull();
    assertThat(response.getCallbackDelaySeconds()).isEqualTo(0);

    assertThat(response.getResourceModel()).isEqualTo(request.getDesiredResourceState());
    assertThat(response.getResourceModels()).isNull();
    assertThat(response.getMessage()).isNull();
    assertThat(response.getErrorCode()).isNull();
}

@Test
public void testHandlerInvokedWhenInstanceIsAlreadyTerminated() {
    final DescribeInstancesResult describeInstancesResult =
        new DescribeInstancesResult().withReservations(new
Reservation().withInstances(new Instance().withState(new
InstanceState().withName("terminated"))

        .withSecurityGroups(new GroupIdentifier().withGroupId("sg-1234"))));

doReturn(describeInstancesResult).when(proxy).injectCredentialsAndInvoke(any(DescribeInstancesRequest.class), any());

    final DeleteHandler handler = new DeleteHandler();

    final ResourceModel model = ResourceModel.builder().instanceId("i-1234").build();

    final ResourceHandlerRequest<ResourceModel> request =
ResourceHandlerRequest.<ResourceModel>builder()
        .desiredResourceState(model)
        .build();

    final ProgressEvent<ResourceModel, CallbackContext> response
        = handler.handleRequest(proxy, request, null, logger);

    assertThat(response).isNotNull();
    assertThat(response.getStatus()).isEqualTo(OperationStatus.FAILED);
    assertThat(response.getCallbackContext()).isNull();
    assertThat(response.getCallbackDelaySeconds()).isEqualTo(0);
    assertThat(response.getResourceModel()).isNull();
    assertThat(response.getResourceModels()).isNull();
    assertThat(response.getMessage()).isNull();
}
```

```

        assertThat(response.getErrorCode()).isEqualTo(HandlerErrorCode.NotFound);
    }

    @Test
    public void testInProgressStateSecurityGroupsNotGathered() {
        final DescribeInstancesResult describeInstancesResult =
            new DescribeInstancesResult().withReservations(new
Reservation().withInstances(new Instance().withState(new
InstanceState().withName("running")))

                .withSecurityGroups(new GroupIdentifier().withGroupId("sg-1234"))));

doReturn(describeInstancesResult).when(proxy).injectCredentialsAndInvoke(any(DescribeInstancesRequest.class),
any());

        final DeleteHandler handler = new DeleteHandler();

        final ResourceModel model = ResourceModel.builder().instanceId("i-1234").build();

        final ResourceHandlerRequest<ResourceModel> request =
ResourceHandlerRequest.<ResourceModel>builder()
            .desiredResourceState(model)
            .build();

        final ProgressEvent<ResourceModel, CallbackContext> response
            = handler.handleRequest(proxy, request, null, logger);

        final CallbackContext desiredOutputContext = CallbackContext.builder()

.stabilizationRetriesRemaining(60)

.instanceSecurityGroups(Arrays.asList("sg-1234"))

                .build();

        assertThat(response).isNotNull();
        assertThat(response.getStatus()).isEqualTo(OperationStatus.IN_PROGRESS);

assertThat(response.getCallbackContext()).isEqualTo(ComparingFieldByFieldComparer.newInstance(desiredOutputContext));
        assertThat(response.getCallbackDelaySeconds()).isEqualTo(0);

assertThat(response.getResourceModel()).isEqualTo(request.getDesiredResourceState());
        assertThat(response.getResourceModels()).isNull();
        assertThat(response.getMessage()).isNull();
        assertThat(response.getErrorCode()).isNull();
    }

    @Test
    public void testInProgressStateSecurityGroupsGathered() {
        final InstanceState inProgressState = new InstanceState().withName("in-
progress");
        final TerminateInstancesResult terminateInstancesResult =
            new TerminateInstancesResult().withTerminatingInstances(new
InstanceStateChange().withCurrentState(inProgressState));

doReturn(terminateInstancesResult).when(proxy).injectCredentialsAndInvoke(any(TerminateInstancesRequest.class),
any());

        final DeleteHandler handler = new DeleteHandler();

        final ResourceModel model = ResourceModel.builder().instanceId("i-1234").build();

        final ResourceHandlerRequest<ResourceModel> request =
ResourceHandlerRequest.<ResourceModel>builder()
            .desiredResourceState(model)
            .build();

        final CallbackContext context = CallbackContext.builder()

```



```

        .stabilizationRetriesRemaining(60)

    .instanceSecurityGroups(Arrays.asList("sg-1234"))
        .build();

    final ProgressEvent<ResourceModel, CallbackContext> response
        = handler.handleRequest(proxy, request, context, logger);

    final CallbackContext desiredOutputContext = CallbackContext.builder()
        .stabilizationRetriesRemaining(60)
        .instanceSecurityGroups(context.getInstanceSecurityGroups())
        .instance(new
    Instance().withState(inProgressState))
        .build();

    assertThat(response).isNotNull();
    assertThat(response.getStatus()).isEqualTo(OperationStatus.IN_PROGRESS);

    assertThat(response.getCallbackContext()).isEqualToComparingFieldByField(desiredOutputContext);
    assertThat(response.getCallbackDelaySeconds()).isEqualTo(0);

    assertThat(response.getResourceModel()).isEqualTo(request.getDesiredResourceState());
    assertThat(response.getResourceModels()).isNull();
    assertThat(response.getMessage()).isNull();
    assertThat(response.getErrorCode()).isNull();
}

@Test
public void testInProgressStateInstanceTerminationInvoked() {
    final InstanceState inProgressState = new InstanceState().withName("in-
progress");
    final GroupIdentifier group = new GroupIdentifier().withGroupId("sg-1234");
    final Instance instance = new
    Instance().withState(inProgressState).withSecurityGroups(group);
    final DescribeInstancesResult describeInstancesResult =
        new DescribeInstancesResult().withReservations(new
    Reservation().withInstances(instance));

    doReturn(describeInstancesResult).when(proxy).injectCredentialsAndInvoke(any(DescribeInstancesRequest.class),
    any());

    final DeleteHandler handler = new DeleteHandler();

    final ResourceModel model = ResourceModel.builder().instanceId("i-1234").build();

    final ResourceHandlerRequest<ResourceModel> request =
    ResourceHandlerRequest.<ResourceModel>builder()
        .desiredResourceState(model)
        .build();

    final CallbackContext context = CallbackContext.builder()
        .stabilizationRetriesRemaining(60)
        .instance(new
    Instance().withState(inProgressState).withSecurityGroups(group))
        .instanceSecurityGroups(Arrays.asList("sg-1234"))
        .build();

    final ProgressEvent<ResourceModel, CallbackContext> response
        = handler.handleRequest(proxy, request, context, logger);

    final CallbackContext desiredOutputContext = CallbackContext.builder()
        .stabilizationRetriesRemaining(59)

```

```
.instanceSecurityGroups(context.getInstanceSecurityGroups())
Instance().withState(inProgressState).withSecurityGroups(group))
                                                    .instance(new
                                                    .build();

    assertThat(response).isNotNull();
    assertThat(response.getStatus()).isEqualTo(OperationStatus.IN_PROGRESS);

assertThat(response.getCallbackContext()).isEqualToComparingFieldByField(desiredOutputContext);
    assertThat(response.getCallbackDelaySeconds()).isEqualTo(0);

assertThat(response.getResourceModel()).isEqualTo(request.getDesiredResourceState());
    assertThat(response.getResourceModels()).isNull();
    assertThat(response.getMessage()).isNull();
    assertThat(response.getErrorCode()).isNull();
}

@Test
public void testStabilizationTimeout() {
    final DeleteHandler handler = new DeleteHandler();

    final ResourceModel model = ResourceModel.builder().instanceId("i-1234").build();

    final ResourceHandlerRequest<ResourceModel> request =
ResourceHandlerRequest.<ResourceModel>builder()
    .desiredResourceState(model)
    .build();

    final CallbackContext context = CallbackContext.builder()
                                                    .stabilizationRetriesRemaining(0)

.instanceSecurityGroups(Arrays.asList("sg-1234"))
                                                    .instance(new
Instance().withState(new InstanceState().withName("terminated")))
                                                    .build();

    try {
        handler.handleRequest(proxy, request, context, logger);
    } catch (RuntimeException e) {
        assertThat(e.getMessage()).isEqualTo(EXPECTED_TIMEOUT_MESSAGE);
    }
}
}
```

Test the Resource Provider

Next, we'll use the AWS SAM CLI to test locally that our resource will work as expected once we submit it to the CloudFormation registry. To do this, we'll need to define tests for the SAM to run against our create and delete handlers.

Create the SAM Test Files

1. Create two files:

- `package-root/sam-tests/create.json`
- `package-root/sam-tests/delete.json`

Where `package-root` is the root of the resource project. For our walkthrough example, the files would be:

- `example-testing-wordpress/sam-tests/create.json`

- `example-testing-wordpress/sam-tests/delete.json`

2. In `example-testing-wordpress/sam-tests/create.json`, paste the following test.

Note

Add the necessary information, such as credentials and log group name, and remove any comments in the file before testing.

To generate temporary credentials, you can use `aws sts get-session-token`.

```
{
  "credentials": {
    # Real STS credentials need to go here.
    "accessKeyId": "",
    "secretAccessKey": "",
    "sessionToken": ""
  },
  "action": "CREATE",
  "request": {
    "clientRequestToken": "4b90a7e4-b790-456b-a937-0cfdfa211dfe", # Can be any UUID.
    "desiredResourceState": {
      "Name": "MyBlog",
      "SubnetId": "subnet-0bc6136e" # This should be a real subnet that exists in
the account you're testing against.
    },
    "logicalResourceIdentifier": "MyResource"
  },
  "callbackContext": null
}
```

3. In `example-testing-wordpress/sam-tests/delete.json`, paste the following test.

Note

Add the necessary information, such as credentials and log group name, and remove any comments in the file before testing.

To generate temporary credentials, you can use `aws sts get-session-token`.

```
{
  "credentials": {
    # Real STS credentials need to go here.
    "accessKeyId": "",
    "secretAccessKey": "",
    "sessionToken": ""
  },
  "action": "DELETE",
  "request": {
    "clientRequestToken": "4b90a7e4-b790-456b-a937-0cfdfa211dfe", # Can be any UUID.
    "desiredResourceState": {
      "Name": "MyBlog",
      "InstanceId": "i-0167b19dd4c1efbf3", # This should be the instance ID that
was created in the "create.json" test.
      "SubnetId": "subnet-0bc6136e" # This should be a real subnet that exists in
the account you're testing against.
    },
    "logicalResourceIdentifier": "MyResource"
  },
  "callbackContext": null
}
```

Test the Create Handler

Once you've created the `example-testing-wordpress/sam-tests/create.json` test file, you can use it to test your create handler.

Ensure Docker is running on your computer.

1. Invoke the SAM function from the resource package root directory using the following commands.

```
sam local invoke TestEntrypoint --event sam-tests/create.json
```

Note

Occasionally these tests will fail with a retry-able error. In such a case, run the tests again to determine whether the issue was transient.

Because the create handler was written as a state machine, invoking the tests will return an output that represents a state. For example:

```
{
  "callbackDelaySeconds": 0,
  "resourceModel": {
    "SubnetId": "subnet-0bc6136e",
    "Name": "MyBlog"
  },
  "callbackContext": {
    "instance": {
      "subnetId": "subnet-0bc6136e",
      "virtualizationType": "hvm",
      "capacityReservationSpecification": {
        "capacityReservationPreference": "open"
      },
      "amiLaunchIndex": 0,
      "elasticInferenceAcceleratorAssociations": [],
      "sourceDestCheck": true,
      "stateReason": {
        "code": "pending",
        "message": "pending"
      },
      "instanceId": "i-0b6978477c0e9d358",
      "vpcId": "vpc-eb80788e",
      "hypervisor": "xen",
      "rootDeviceName": "/dev/sda1",
      "productCodes": [],
      "state": {
        "code": 0,
        "name": "pending"
      },
      "architecture": "x86_64",
      "ebsOptimized": false,
      "imageId": "ami-04fb0368671b6f138",
      "blockDeviceMappings": [],
      "stateTransitionReason": "",
      "clientToken": "207dc686-e95c-4df9-8fcb-ee22bbdde963",
      "instanceType": "m4.large",
      "cpuOptions": {
        "threadsPerCore": 2,
        "coreCount": 1
      },
      "monitoring": {
        "state": "disabled"
      },
      "publicDnsName": "",
      "privateIpAddress": "172.0.0.133",
      "rootDeviceType": "ebs",
      "tags": [
        {
          "value": "MyBlog",
          "key": "Name"
        }
      ]
    }
  }
}
```

```
    }
  ],
  "launchTime": 1567718644000,
  "elasticGpuAssociations": [],
  "licenses": [],
  "networkInterfaces": [
    {
      "networkInterfaceId": "eni-0e450b35a159b60fe",
      "privateIpAddresses": [
        {
          "privateIpAddress": "172.0.0.133",
          "primary": true
        }
      ],
      "subnetId": "subnet-0bc6136e",
      "description": "",
      "groups": [
        {
          "groupName": "MyBlog-cbb70fca-4704-430b-b67b-7d6d550e0592",
          "groupId": "sg-063679dc7681610c3"
        }
      ],
      "ipv6Addresses": [],
      "ownerId": "671472782477",
      "sourceDestCheck": true,
      "privateIpAddress": "172.0.0.133",
      "interfaceType": "interface",
      "macAddress": "02:e1:4b:d1:f7:40",
      "attachment": {
        "attachmentId": "eni-attach-0a01c63e4b45c4a5d",
        "deleteOnTermination": true,
        "deviceIndex": 0,
        "attachTime": 1567718644000,
        "status": "attaching"
      },
      "vpcId": "vpc-eb80788e",
      "status": "in-use"
    }
  ],
  "privateDnsName": "ip-172-0-0-133.us-west-2.compute.internal",
  "securityGroups": [
    {
      "groupName": "MyBlog-cbb70fca-4704-430b-b67b-7d6d550e0592",
      "groupId": "sg-063679dc7681610c3"
    }
  ],
  "placement": {
    "groupName": "",
    "tenancy": "default",
    "availabilityZone": "us-west-2b"
  },
  "stabilizationRetriesRemaining": 60
},
"status": "IN_PROGRESS"
}
```

2. From the test response, copy the contents of the callbackContext, and paste it into the callbackContext section of the example-testing-wordpress/sam-tests/create.json file.
3. Invoke the TestEntrypoint function again.

```
sam local invoke TestEntrypoint --event sam-tests/create.json
```

If the resource has yet to complete provisioning, the test returns a response with a status of `IN_PROGRESS`. Once the resource has completed provisioning, the test returns a response with a status of `SUCCESS`. For example:

```
{
  "callbackDelaySeconds": 0,
  "resourceModel": {
    "InstanceId": "i-0b6978477c0e9d358",
    "PublicIp": "34.211.69.121",
    "SubnetId": "subnet-0bc6136e",
    "Name": "MyBlog"
  },
  "status": "SUCCESS"
}
```

4. Repeat the previous two steps until the resource has completed.

When the resource completes provisioning, the test response contains both its `PublicIp` and `InstanceId`:

- You can use the `PublicIp` value to navigate to the WordPress site.
- You can use the `InstanceId` value to test the delete handler, as described below.

Test the Delete Handler

Once you've created the `example-testing-wordpress/sam-tests/delete.json` test file, you can use it to test your delete handler.

Ensure Docker is running on your computer.

1. Invoke the `TestEntrypoint` function from the resource package root directory using the following commands.

```
sam local invoke TestEntrypoint --event sam-tests/delete.json
```

Note

Occasionally these tests will fail with a retry-able error. In such a case, run the tests again to determine whether the issue was transient.

As with the create handler, the delete handler was written as a state machine, so invoking the test will return an output that represents a state.

2. From the test response, copy the contents of the `callbackContext`, and paste it into the `callbackContext` section of the `example-testing-wordpress/sam-tests/delete.json` file.
3. Invoke the `TestEntrypoint` function again.

```
sam local invoke TestEntrypoint --event sam-tests/delete.json
```

If the resource has yet to complete provisioning, the test returns a response with a status of `IN_PROGRESS`. Once the resource has completed provisioning, the test returns a response with a status of `SUCCESS`.

4. Repeat the previous two steps until the resource has completed.

Performing Resource Contract Tests

Resource contract tests verify that the resource type provider schema you've defined properly catches property values that will fail when passed to the underlying APIs called from within your resource handlers. This provides a way of validating user input before passing it to the resource handlers. For example, in the `Example::Testing::WordPress` resource type provider schema (in the `example-testing-wordpress.json` file), we specified regex patterns for the `Name` and `SubnetId` properties, and set the maximum length of `Name` as 219 characters. Contract tests are intended to stress and validate those input definitions.

Specify Resource Contract Test Override Values

The CloudFormation CLI performs resource contract tests using input that is generated from the patterns you define in your resource's property definitions. However, some inputs can't be randomly generated. For example, the `Example::Testing::WordPress` resource requires an actual subnet ID for testing, not just a string that matches the appearance of a subnet ID. So in order to test this property, we need to include a file with actual values for the resource contract tests to use. an `overrides.json` at the root of our project that looks like this:

1. Navigate to the root of your project.
2. Create a file named `overrides.json`.
3. Include the following override, specifying an actual subnet ID to use when performing resource contract tests.

```
{
  "CREATE": {
    "/SubnetId": "subnet-0bc6136e" # This should be a real subnet that exists in the
    account you're testing against.
  }
}
```

Run the Resource Contract Tests

To run resource contract tests, you'll need two shell sessions.

1. In a new session, run the following command:

```
$ sam local start-lambda
```

2. From the resource package root directory, in a session that is aware of the CloudFormation CLI, run the test command:

```
$ cfn test
```

The session that is running `sam local start-lambda` will display information about the status of your tests.

Submit the Resource Provider

Once you have completed implementing and testing your resource provided, the final step is to submit it to the CloudFormation registry. This makes it available for use in stack operations in the account and region in which it was submitted.

- In a terminal, run the `submit` command to register the resource provider in the `us-west-2` region.

```
$ cfn submit -v --region us-west-2
```

The CloudFormation CLI validates the included resource provider schema, packages your resource provide project and uploads it to the CloudFormation registry, and then returns a registration token.

```
Validating your resource specification...
Packaging Java project
Creating managed upload infrastructure stack
Managed upload infrastructure stack was successfully created
Registration in progress with token: 3c27b9e6-dca4-4892-ba4e-3c0example
```

Resource provider registration is an asynchronous operation. You can use the supplied registration token to track the progress of your provider registration request using the [DescribeTypeRegistration](#) action of the CloudFormation API.

Note

If you update your resource provider, you can submit a new version of that resource provider. Every time you submit your resource provider, CloudFormation generates a new version of that resource provider.

To set the default version of a resource provider, use [SetTypeDefaultVersion](#). For example:

```
aws cloudformation set-type-default-version --type "RESOURCE" --type-name
"Example::Testing::WordPress" --version-id "00000002"
```

To retrieve information about the versions of a resource provider, use [ListTypeVersions](#). For example:

```
aws cloudformation list-type-versions --type "RESOURCE" --type-name
"Example::Testing::WordPress"
```

Provision the Resource in a CloudFormation Stack

Once the registration request for your resource provider has completed successfully, you can create a stack including resources of that type.

Note

Use [DescribeTypeRegistration](#) to determine when your resource provider is successfully registration registered with a status of `COMPLETE`. You should also see your new resource provider listed in the [CloudFormation console](#).

1. Save the following JSON as a stack template, with the name `stack.json`.

```
{
  "AWSTemplateFormatVersion": "2010-09-09",
  "Description": "WordPress stack",
  "Resources": {
    "MyWordPressSite": {
      "Type": "Example::Testing::WordPress",
      "Properties": {
        "SubnetId": "subnet-0bc6136e", ## Note that this should be replaced with
a subnet that exists in your account.
        "Name": "MyWebsite"
      }
    }
  }
}
```



```
}
```

2. Use the template to create a stack.

Note

This resource uses an official WordPress image on AWS Marketplace. In order to create the stack, you'll first need to visit the [AWS Marketplace](#) and accept the terms and subscribe.

Navigate to the folder in which you saved the `stack.json` file, and create a stack named `wordpress`.

```
aws cloudformation create-stack --region us-west-2 \  
--template-body "file://stack.json" \  
--stack-name "wordpress"
```

As CloudFormation creates the stack, it should invoke your resource provider create handler to provision a resource of type `Example::Testing::WordPress` as part of the `wordpress` stack.

As a final test of the resource provider delete handler, you can delete the `wordpress` stack.

```
aws cloudformation delete-stack --region us-west-2 \  
--stack-name wordpress
```

Resource Provider FAQ

Below are some Frequently Asked Questions about resource provider development.

Updates

- **Q: My service API implements Update actions as an Upsert, can I implement my CloudFormation resource provider in this way?**

A: No, CloudFormation requires that update actions to a non-existing resource always throw a `ResourceNotFoundException`.

Schema Development

- **Q: How can I re-use existing schemas, or establish relationships to other resource types in my schema?**

A: Relationships and re-use are established using JSON Pointers. These are implemented using the `$ref` keyword in your resource provider schema. Refer to [Modeling Resource Providers for Use in AWS CloudFormation \(p. 4\)](#) for more information.

Permissions and Authorization

- **Q: Why am I getting an `AccessDeniedException` for my AWS API calls?**

A: If you are seeing errors in your logs related to `AccessDeniedException` for a Lambda Execution Role like

```
A downstream service error occurred in a CREATE action on a  
AWS::MyService::MyResource: com.amazonaws.services.logs.model.AWSLogsException: User:
```

```
arn:aws:sts::123456789012:assumed-role/  
UluruResourceHandlerLambdaExecutionRole-123456789012pdx/AWS-MYService-MyResource-  
Handler-1h344teffe is not authorized to perform: some:ApiCall on resource: some-resource  
(Service: AWSLogs; Status Code: 400; Error Code: AccessDeniedException; Request ID:  
36af0cec-a96a-11e9-b204-ddabexample)
```

This is an indication that you are attempting to create and invoke AWS APIs using the default client, which is injected with environment credentials.

For resource providers, you should use the passed-in `AmazonWebServicesClientProxy` object to make AWS API calls, as in the following example.

```
SesClient client = ClientBuilder.getClient();  
final CreateConfigurationSetRequest createConfigurationSetRequest =  
    CreateConfigurationSetRequest.builder()  
        .configurationSet(ConfigurationSet.builder()  
            .name(model.getName())  
            .build())  
        .build();  
proxy.injectCredentialsAndInvokeV2(createConfigurationSetRequest,  
    client::createConfigurationSet);
```

- **Q: How do I specify credentials for non-AWS API calls?**

A: For non-AWS API calls which require authentication and authorization, you should create properties in your resource provider which contain the credentials. Define these properties in the resource provider schema as `writeOnlyProperties`.

Users can then provide their own credentials through their CloudFormation templates. We encourage the use of [dynamic references](#) in CloudFormation templates, which can use AWS Secrets Manager to fetch credentials at runtime.

Provider Development

- **Q: Can I share functionality between resource by adding common functionality to the `BaseHandler`?**

A: Because the `BaseHandler` is code-generated, it cannot be edited.

- **Q: For Java development, is there a way to include multiple resources in a single maven project?**

A: Not currently. For security and manageability, the CloudFormation Registry registers each resource provider as a separate, versioned, type. You could still share code through a shared package. Ideally, the wrapper layer does most of the boilerplate. If you see a need for more boilerplate, we would like to know how we can improve for that use case rather than combine types in a package, so please reach out to the team.

- **Q: Will `software.amazon.cloudformation.proxy.Logger` have debug/info/warning/error levels/?**

A: Currently, all log messages are emitted to AWS CloudWatch, which has no built-in concept of log levels.

Testing

- **Q: For testing, when should I use `sam local invoke`, `cfn test` and `mvn test`?**

A: Use the various test capabilities for the test scenarios described below.

- `sam local invoke`: Creating custom integration test by passing in custom CloudFormation payloads and isolate specific handlers.
- `cfn test`: Contract tests meant to cycle through CRUDL actions and ideally leave in a clean state (if tests pass).
- `mvn test`: Used for Unit testing. The goal is to confirm that each method/unit of the resource performs as intended. It also helps to ensure that you have enough code coverage. We expect unit tests to mock dependencies and not create real resources.

Deployment

- **Q: Is the provider interface guaranteed to be stable?**

A: The communication protocol between CloudFormation and your provider resource package is subject to change. However this will be done in a backwards-compatible way, using versioned interfaces. This will be invisible to you as a developer and is managed as part of the CloudFormation managed platform.

The interface that your handlers implement inside your package is expected to be stable. We may introduce improvements, such as security fixes or other changes to the package, but these will be done with versioned dependency or CloudFormation CLI updates. You are not required to upgrade your packages in order to publish them, only to incorporate these improvements.

Document History for User Guide

The following table describes the documentation for this release of the CloudFormation Command Line Interface (CLI).

- **Latest documentation update:** November 14, 2019

update-history-change	update-history-description	update-history-date
General Availability (p. 65)	Initial publication of the CloudFormation CLI documentation.	November 14, 2019

AWS glossary

For the latest AWS terminology, see the [AWS glossary](#) in the *AWS General Reference*.